

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/340042170>

Geometric Deep Learning on Skeleton Sequences for 2D/3D Action Recognition

Conference Paper · January 2020

DOI: 10.5220/0009161701960204

CITATIONS

2

READS

252

3 authors:



Rasha Friji

Cristal Lab- Université de la Manouba

3 PUBLICATIONS 2 CITATIONS

[SEE PROFILE](#)



Hassen Drira

Institut Mines-Télécom

49 PUBLICATIONS 798 CITATIONS

[SEE PROFILE](#)



Faten Chaieb

Ecole Nationale des Sciences de l'Informatique

39 PUBLICATIONS 99 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



TAIMA - Treatment and Analysis of Information Methods and Applications. [View project](#)



Image registration [View project](#)

Geometric Deep Learning on Skeleton Sequences for 2D/3D Action Recognition

Rasha Friji^{1,2}, Hassen Drira³ and Faten Chaieb⁴

¹*CRISTAL Lab, National University of Computer Science ENSI, Manouba University Campus, Manouba, Tunisia*

²*Talan Innovation Factory, Talan, Tunisia*

³*IMT Lille Douai, Univ. Lille, CNRS, UMR 9189,*

CRISTAL – Centre de Recherche en Informatique Signal et Automatique de Lille, F-59000 Lille, France

⁴*Ecole Nationale des Sciences de l'Informatique INSAT, Tunisia*

<https://talan.com/>

Keywords: Geometric Deep Learning, Action Recognition, Abnormal Gait Recognition.

Abstract: Deep Learning models, albeit successful on data defined on Euclidean domains, are so far constrained in many fields requiring data which underlying structure is a non-Euclidean space, namely computer vision and imaging. The purpose of this paper is to build a geometry aware deep learning architecture for skeleton based action recognition. In this perspective, we propose a framework for non-Euclidean data classification based on 2D/3D skeleton sequences, specifically for Parkinson's disease classification and action recognition. As a baseline, we first design two Euclidean deep learning architectures without considering the Riemannian structure of the data. Then, we introduce new architectures that extend Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs) to non-Euclidean data. Experimental results show that our method outperforms state-of-the-art performances for 2D abnormal behavior classification and 3D human action recognition.

1 INTRODUCTION

Geometric deep learning is a terminology, initiated by Bronstein et al. (Bronstein et al., 2017), and used to refer to deep learning approaches to generalize deep neural networks to non-Euclidean domains such as manifolds (e.g action recognition) and graphs (image analysis). Skeleton sequences is an example of non-Euclidean data which has been increasingly standing out (Du et al., 2015; Shahroudy et al., 2016; Vemulapalli et al., 2014; Ke et al., 2017) given the availability of huge datasets and the multitude of possible applications. In this paper, we focus on skeleton-based 2D Parkinson's disease classification and 3D skeleton-based action recognition. CNNs (LeCun and Bengio, 1998) have proven distinguished performance in image classification (Krizhevsky et al., 2012; Cohen et al., 2018; Xiong et al., 2015; Ke and Li, 2014; Ciresan et al., 2012). In our work, we propose a first CNN based architecture validated for the Parkinson's disease classification. But, instead of directly applying CNN on video images, we primarily identify the sequence of performed actions, by tracking the trajectories of human skeleton joints. We represent

then every action with a sequence of joints' 2D coordinates and perform spherical modelling followed by a projection in the tangent space. Nonetheless, as far as video action recognition is concerned, even deep CNNs are still not capable of modelling the temporal correlation between the video frames (Wang et al., 2016). In order to address this limitation and to exploit the dynamics of human movements, the joints' series have been used in recurrent neural networks (RNNs) with Long-Short Term Memory (LSTM) neurons (Graves, 2012; Graves et al., 2013) for action recognition (Du et al., 2015; Veeriah et al., 2015; Zhu et al., 2016). Motivated by these works, we propose a second non-Euclidean architecture based on CNN and LSTM combination for action recognition, tested on NTU RGB+D dataset. At the first layer of the network, we used a deep CNN for features extraction. At the second layer, features are then passed to LSTM which makes the network temporally-aware. **The contributions of this work are** 1) Novel non-Euclidean deep neural networks architectures for 2D/3D skeletal sequences based action recognition. 2) Actions are recognized with respect to the geometry of the manifold of skeletal sequences and with

respect to the temporal dependencies between these sequences 3) Ablation studies for the classification of Parkinson's disease and action recognition. Experiments are conducted on two benchmark datasets to prove the competitiveness of the proposed method.

The rest of the paper is organized as follows. In section 2, we briefly review existing geometry-aware deep learning models applied on non linear manifolds using CNNs and RNNs networks. Section 3 introduces the spherical modelling of skeletal data and the mapping to the tangent space. In section 4, we describe the proposed method. Experimental settings and results are reported in section 5, and lastly section 6 concludes the paper.

2 RELATED WORK

In this section, we briefly review the relevant literature of geometry-aware deep learning architectures for skeleton sequences based classification using CNNs and RNNs networks.

2.1 CNNs based Methods

Unlike the common deep learning architectures that have been widely used in many applications, only limited efforts have been spent on non-linear deep learning. In the past few years, the interest for CNNs adapted to manifolds has grown exponentially. The first generalized CNNs to manifolds was proposed by (LeCun and Bengio, 1998) who used local intrinsic patches to define the convolution operation. Yann LeCun and M.Bronstein proposed in (Bronstein et al., 2017) an overview of the mainly used non-Euclidean deep learning architectures. In problems like shape description, retrieval, and correspondence, a Geodesic Convolutional Neural Network (GCNN) was designed in (Masci et al., 2015) to learn invariant shape features. GCNN is an intrinsic version of CNN on manifolds where Masci, Jonathan, et al proposed an application of filters on local patches represented in geodesic polar coordinates. In (Cohen et al., 2018), Cohen, Taco S., et al., proposed spherical CNNs using a generalized Fast Fourier Transform (FFT).

2.2 CNN and RNNs Combination based Methods

Recently, RNNs (Baccouche et al., 2011; Lefebvre et al., 2013; Zhu et al., 2016) have been used for action recognition. However, vanishing gradient problems often occurred because of the large number of

parameters computations and the neglect of initial input effect after few layers. As a solution, LSTM networks (Donahue et al., 2017a; Ng et al., 2015; Srivastava et al., 2015) were deployed since they integrate memory units and they are subsequently capable of capturing long-term dependencies. Based on the extension of CNNs to 3D, Baccouche et al. (Baccouche et al., 2011) propose a unidirectional model with only one hidden layer LSTM-RNN for action recognition. Lefebvre et al. (Lefebvre et al., 2013) propose a bidirectional LSTM-RNN with one forward hidden layer and one backward hidden layer for gesture classification. In (Zhu et al., 2016), in order to learn the inherent co-occurrence features of skeleton joints, these joints are fed to a regularized deep LSTM at each time interval. In (Shahroudy et al., 2016), Shahroudy et al. propose to learn the long-term context representations of the body parts with a part-aware LSTM (P-LSTM). In (Liu et al., 2016), a spatial temporal LSTM is used to learn both the spatial and temporal information of skeleton sequences and a Trust Gate is introduced to omit noisy joints. This approach achieves the state-of-the-art performance on the NTU RGB+D dataset (Shahroudy et al., 2016).

3 MODELLING OF 2D/3D SKELETAL DATA

Actions captured by visual sensors and cameras are often subject to scale variations. This is due to the change of distance between the camera and the person performing the action. As a result, the same actions can be interpreted very differently. In order to avoid this problem, sequences of skeletons should be invariant to global scaling. For this purpose, a modeling is done on the input of our architecture to normalize skeletons.

3.1 Spherical Modelling

Let $X \in \mathcal{R}^{n \times k}$ be a body skeleton, where n indicates the number of body joints and k denotes the dimension of X . To remove scale, we propose to model skeletons as elements on a $(n \times k - 1)$ dimension Riemannian manifold, more specifically, the unit sphere S embedded in $\mathcal{R}^{n \times k}$. To do so, we divide every skeleton X by its Frobenius norm given by Eq.1:

$$\|X\|_F = \left(\sum_{i,j=1}^n |x_{ij}|^2 \right)^{1/2} \quad (1)$$

With this process, we consequently get skeletons representations as well as their temporal evolution,

called trajectories on the unit sphere S embedded in $\mathcal{R}^{n \times k}$. Accordingly, each motion sequence of a moving skeleton is represented with a trajectory on the unit sphere S embedded in $\mathcal{R}^{n \times k}$ as shown in Fig.1 (for visualization purposes the trajectory is shown in 2D, however, it lies, in fact, on a space of $(n \times k - 1)$ dimensions).




Figure 1: Spherical modelling of the skeleton sequence data.

3.2 Inverse Exponential Map

The input fed to our CNNs network lies on a Riemannian manifold which is the Unit Sphere S embedded in $\mathcal{R}^{n \times k}$. Unlike the Euclidean space which is a vector space characterized by translation invariance and operations like vector addition and scalar multiplication, the non-Euclidean structure of our input implies that there are no such properties. Consequently, even basic operations like convolution can't be applied on the Unit Sphere S embedded in $\mathcal{R}^{n \times k}$, since they are not defined. However, manifolds, including Riemannian manifold, are topological spaces that can be locally assimilated to an Euclidean space. Given that the unit sphere S embedded in $\mathcal{R}^{n \times k}$ has a Riemannian manifold structure, the manifold can be assimilated, locally around each point x , to an Euclidean space known as the tangent space $T_X(S)$.

Following, we define the tangent space shown in Fig.2 and the inverse exponential map layer used to map data from the Riemannian manifold which is the unit sphere embedded in $\mathcal{R}^{n \times k}$ to a tangent space.

A differentiable d -dimensional manifold X is a topological space where each point x has a neighborhood, which is homeomorphic to a d -dimensional Euclidean space, a.k.a the tangent space and denoted by $T_x(X)$. In other words, at each point x on the manifold X , it is possible to associate a linear space $T_x(X)$. The space $T_x(X)$ is a local Euclidean representation of the manifold X around x . This space is called the tangent space of the manifold X at the point x . Considering that the tangent space is linear and hence equipped




Figure 2: Examples of two tangent spaces: $T_x(X)$ at a point x of the manifold X and $T_{x'}(X)$ at a point x' of the manifold X .

with the inner product, the Riemannian metric on S is defined by Eq.2:

$$\langle X_1, X_2 \rangle = \text{trace}(X_1, X_2), X_1, X_2 \in T_X(S) \quad (2)$$




Figure 3: Unit Sphere S embedded in $\mathcal{R}^{n \times k}$, the trajectories α_1 and α_2 of two sequences of skeletons, the geodesic $\alpha(t)$ connecting arbitrary points on α_1 and α_2 , the tangent space $T_{X_1}(S)$ at the skeleton X_1 and skeletons X_2 and X_3 mapped on $T_{X_1}(S)$.

The inverse exponential map shown in Fig.3, also known as the logarithm map and uniquely defined around a small neighborhood of a point x on the manifold X , is given by Eq.3:

$$\exp_{X_1}^{-1}(X_j) = \frac{\theta}{\sin \theta} (X_j - \cos(\theta) X_1) \quad (3)$$

With $\theta = \cos^{-1}(\text{trace}(X_1(X_j)^T))$. Here X_1 and X_j represent skeletons on the unit S embedded in $\mathcal{R}^{n \times k}$.

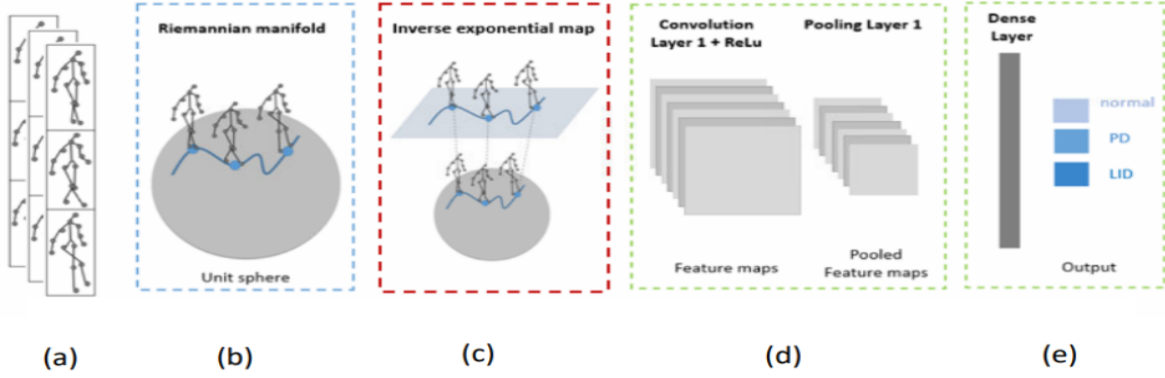


Figure 4: Architecture of the non-Euclidean CNN based proposed method: a) Input 2D skeletal joints coordinates b) Modelling of skeletal data on the Riemannian manifold c) Skeletal data mapping on tangent space d) Feature Extraction with CNN e) Classification.

4 PROPOSED METHOD

Overall architecture of the proposed methods are depicted in Fig.4 and Fig.5. Both models have similar global structure components with the difference that in the first architecture, feature extraction is performed with CNN only, while the second architecture is based on CNN-LSTM, taking into account sequence dependencies.

4.1 Non-Euclidean CNN based Model

In this section, we present a general framework allowing to design CNN architectures on non-Euclidean domains. For this purpose, we build a network structure where each input is an element of the unit sphere S embedded in $\mathcal{R}^{n \times k}$. As shown in Fig.5, this architecture is composed of classical convolution layer, pooling layer, and fully connected layer, subsequent to a spherical modelling and an inverse exponential map layer to address the problem of the non-Euclidean structure of the input data.

After spherical modelling of the skeleton sequence data, we use Eq.3 to map each skeleton X_j from the sphere S to the tangent space $T_{X_1}(S)$ at the skeleton X_1 . We choose a skeleton X_1 as a reference and map all the other skeletons to the tangent space of X_1 as shown in Fig.2. Since the tangent space is an Euclidean space, the input is no more a trajectory on a manifold. It lies however on an Euclidean Space and hence can be fed into any regular CNN layer.

4.2 Non-Euclidean CNN-LSTM based Model

Fig.5 depicts the non Euclidean CNN-LSTM proposed method. This architecture is an extension of

the previous one, aiming to improve and consolidate the obtained results using a better performing model and tested on a larger dataset. As far as the overall building components are concerned, the two architectures are basically identical: input skeletal joints coordinates, spherical modelling, mapping to tangent space, feature extraction and action recognition. The only difference is the introduction of LSTM to capture global sequence dependencies of the input data. CNN-LSTM (Donahue et al., 2017b) architecture involves using CNN layers for feature extraction from input data combined with LSTMs for sequential features interpretation. In our approach, we implement this architecture using two consecutive CNN layers ahead of dropout and a max pooling layer. The whole CNN model is wrapped in a "TimeDistributed" layer. The extracted features are next flattened and provided to the LSTM model before a dense mapping to an action is performed.

5 EXPERIMENTS

In this section, we introduce the testing datasets, the evaluation protocols and the experimental results obtained by our methods with comparison to state of the art and baseline models.

5.1 Datasets

The first proposed architecture has been tested on Parkinson's Vision-Based Pose Estimation Dataset (Li et al., 2017). The second architecture has been tested on NTU RGB+D dataset. In this part, we introduce the two datasets on which we performed our experiments.

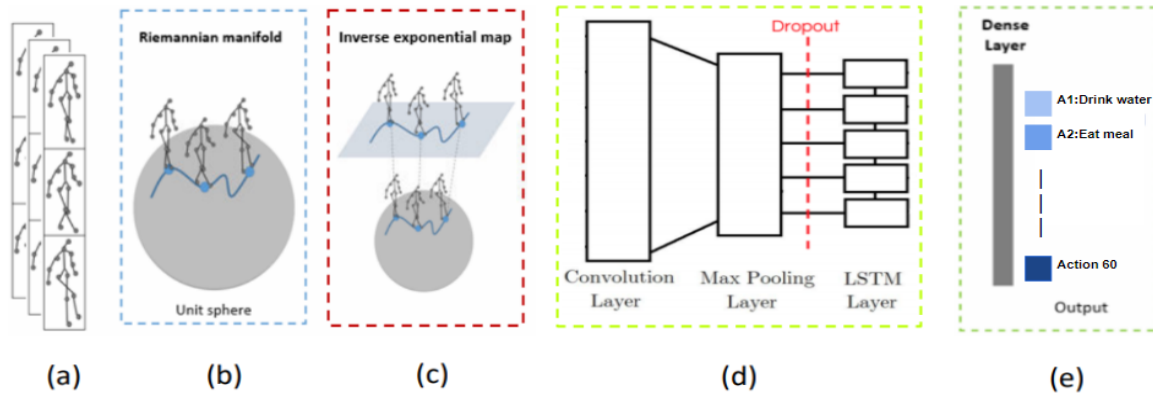


Figure 5: Architecture of the non-Euclidean CNN-LSTM based proposed method: a) Input 3D skeletal joints coordinates b) Modelling of skeletal data on the Riemannian manifold c) Skeletal data mapping on tangent space d) Feature Extraction with CNN combined with LSTM neurons to support sequence prediction e) Action recognition.

Parkinson’s Dataset. Parkinson’s Disease (PD) is a progressive neurodegenerative disease (Nussbaum and Ellis, 2003). It causes a decrease of dopamine neurons in the brain and therefore a reduction of dopamine levels that these neurons produce. The “Parkinson’s Vision-Based Pose Estimation Dataset” (Li et al., 2017) contains movement trajectories extracted from 526 videos of 9 participants (5 men and 4 women) with an average age of 64 years and diagnosed with idiopathic Parkinson’s diseases. Participants completed a two hour Levodopa infusion protocol followed by a two hour observation during which they were asked to do various tasks like drinking from a cup, communication tasks (describing an image, talking with another person, recalling something and mental math), leg agility (stomping of the leg vertically with the maximum possible amount of speed and amplitude) and toe-tapping. Communication and drinking from a cup tasks were used to evaluate the Dyskinesia whereas leg agility and toe-tapping were used to evaluate Parkinsonism. Videos were recorded using a consumer grade video camera and then a 2D human pose estimation was done using Convolutional Pose Machines (CPM) (Wei et al., 2016). Since the CPM algorithm gives only an annotation of the head which is not adapted for tracking head turning movements, an object tracker algorithm was used to estimate the face position. Finally, skeletons with 15 joints are obtained. Since this dataset has action-sequences of variable lengths, we split every sequence into 100-frame sequences, which makes, in total, a number of 30859 sequences. After data sampling, every sequence is composed of 100 skeletons represented with 15 joints. Every joint has two coordinates. The sequences used have therefore a dimension of 100×30 each: 100 indicates the number of skeletons in a sequence and 30 denotes the number of joint-coordinates of each skeleton. For the designed

CNN architecture, we transform those sequences into a 1D signal by concatenating all the skeletons in a sequence.

NTU RGB+D Dataset. This dataset (Shahroudy et al., 2016) is one of the largest skeleton-based human action datasets, consisting of more than 56000 sequences and 4 million frames. It covers 60 classes of actions performed by 40 distinct subjects, including both individual daily actions (e.g. reading, clapping, writing, sneezing, staggering, falling down, etc) and interaction actions (e.g. hugging, handshaking, pointing).

- **Data Modalities:** To collect this dataset, Microsoft Kinect v2 sensors were utilized. Four major data modalities were collected: depth maps, 3D joint information, RGB frames, and IR sequences. Joint information, which is the data modality used in this work, consists of 3-dimensional locations of 25 major body joints for detected and tracked human bodies in the scene. The configuration of body joints is illustrated in Fig.6.

- **Views:** Three cameras were used at the same time to capture three different horizontal views from the same action. For each setup, the three cameras were located at the same height but from three different horizontal angles: -45° , 0° , $+45^\circ$. Each subject was asked to perform each action twice, once towards the left camera and once towards the right camera. Hence, two front views, one left side view, one right side view, one left side 45 degrees view, and one right side 45 degrees view are captured. The three cameras are assigned consistent camera numbers. Camera 1 always observes the 45 degrees views, while camera 2 and 3 observe front and side views.

NTU RGB+D dataset is considered very challenging given the large view points, the intra-class non uniformity and the variation of sequence length. To




Figure 6: Configuration of 25 body joints in our dataset. The labels of the joints are: 1-base of the spine 2-middle of the spine 3-neck 4-head 5-left shoulder 6-left elbow 7-left wrist 8-left hand 9-right shoulder 10-right elbow 11-right wrist 12-right hand 13-left hip 14-left knee 15-left ankle 16-left foot 17-right hip 18-right knee 19-right ankle 20-right foot 21-spine 22-tip of the left hand 23-left thumb 24-tip of the right hand 25-right thumb. (Shahrourdy et al., 2016).

overlap the complication of this variation, we consider the same length of sequence for all the subjects, which is the length of the shortest sequence.

5.2 Evaluation Protocols

As a means to standardize the assessment of the results of our work, we define in this section, the adopted evaluation setups for each of the two datasets.

5.2.1 Protocol on Parkinson’s dataset

Our goal is to classify sequences into three different classes: normal, Parkinson’s disease (PD) or Parkinson’s disease with Levodopa-induced dyskinesia (PD with LID). Given that only the communication sequences have ratings for PD and LID and that the communication task had the best performance according to (Li et al., 2017), we only use those sequences for our multi-classification problem. In all experiments, we adopt the leave-one-out cross validation protocol which means using, for N times (N is the number of instances), the sequences of one person as the validation set and the remaining sequences as the training set.

5.2.2 Protocols for NTU RGB+D

To get standard evaluations for all the reported results on this benchmark, we define precise criteria for two types of action classification evaluation, as described in this section.

Cross-subject Protocol. For the cross-subject evaluation protocol, we split the 40 subjects into training and testing sets, each is composed of 20 subjects. The training and testing groups are made up of 40,320 and 16,560 samples, respectively. In our work, we use for training, the subjects which IDs are among the following list of values: 1, 2, 4, 5, 8, 9, 13, 14, 15, 16, 17, 18, 19, 25, 27, 28, 31, 34, 35, 38. The 20 remaining subjects are reserved for testing.

Cross-view Protocol. In cross-view protocol, we select the samples from cameras 2 and 3 for training and the samples from camera 1 for testing. The training set consists then of the front and two side views of the actions, whilst testing set incorporates left and right 45 degree views of the action performances. For this assessment, the training and testing sets have 37,920 and 18,960 samples, respectively.

5.3 Results

This section summarizes the experimental results for both non Euclidean architectures, each tested on a separate benchmark dataset. The classification accuracy reported in the results sections is in percentage.

5.3.1 Results of CNN Only based Architecture

Table.1 reports the resulting accuracy values of 2D skeletal sequences classification on “Parkinson’s Vision-Based Pose Estimation Dataset” for different methods. In line with the state of the art approach (Li et al., 2017) based on random forest tree algorithm, we compare results achieved using only the communication sequences and adopting leave-one-out cross validation. The Euclidean CNN architecture shown

Table 1: Results on Parkinson’s Vision-Based Pose Estimation Dataset.

Method	Leave-one-out cross validation Accuracy
State of the art (Li et al., 2017)	71.4%
Euclidean CNN (Baseline)	68%
Non Euclidean CNN	72%

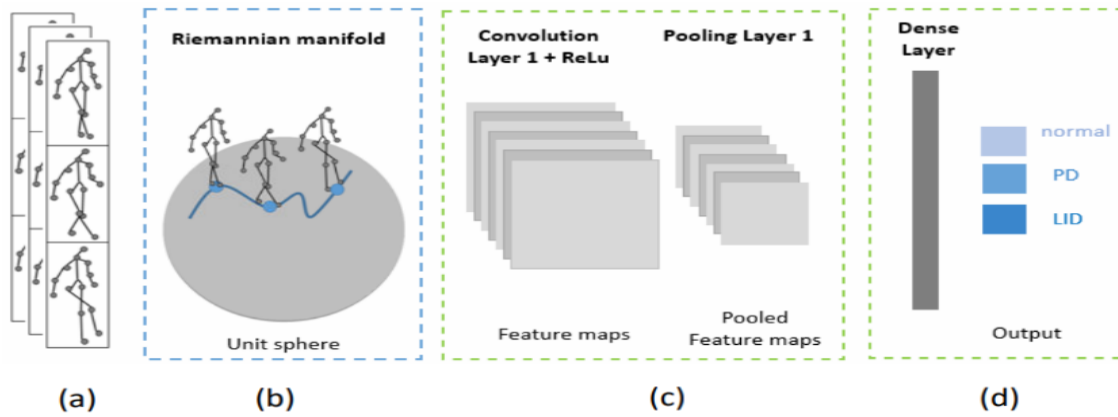


Figure 7: Architecture of the baseline Euclidean CNN based proposed method: a) Input 2D skeletal joints coordinates b) Modelling of skeletal data on the Riemannian manifold c) Feature Extraction with CNN d) Classification.

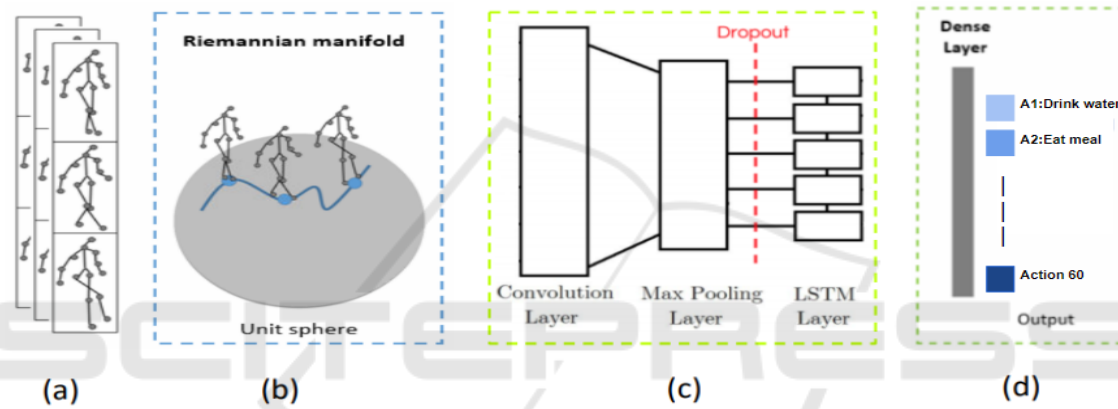


Figure 8: Architecture of the baseline Euclidean CNN-LSTM based proposed method: a) Input 3D skeletal joints coordinates b) Modelling of skeletal data on the Riemannian manifold c) Feature Extraction with CNN combined with LSTM neurons to support sequence prediction d) Action recognition.

in Fig.7 is used as a baseline method to point out the contribution of inverse exponential map layer.

In Table.1, it can be seen that the non-Euclidean CNN architecture, with 72% accuracy, outperforms the baseline architecture which accuracy is 68%. This improvement can highlight the importance of the mapping of the skeletal data to tangent space. Compared with the state of the art results, the performance of our proposed method improved by 0.6%. This improved performance is due to the fact that our method takes into account the non-Euclidean structure of the skeletal data.

5.3.2 Results of CNN-LSTM based Architecture

The results of action recognition on NTU RGB+D dataset for the two evaluation protocols: cross-subject and cross-view are reported in Table.2. The first two lines refer to the accuracy values obtained by the two state of the art approaches(Shahroudy et al., 2016) which deploy respectively, one layer and two layers

Table 2: Results on NTU RGB+D dataset using two evaluation protocols: cross-subject and cross-view. The two first lines refer to the state of the art results(Shahroudy et al., 2016).

Method	Cross-Subject Accuracy	Cross-View Accuracy
1 Layer P-LSTM(Shahroudy et al., 2016)	62.05%	69.40%
2 Layer P-LSTM(Shahroudy et al., 2016)	62.93%	70.27%
Euclidean CNN-LSTM (Baseline)	56.61%	62.32%
Non Euclidean CNN-LSTM	61.45%	71.03%

of part-aware extension of the long short-term memory (P-LSTM). Similarly to CNN based architecture, we used a baseline euclidean CNN-LSTM architec-

ture, as shown in Fig.7, for the purpose of pointing out the contribution of inverse exponential map layer. With reference to this architecture, the performance of our proposed method improves with 4.84% using cross-subject protocol and with 8.71% using cross-view protocol.

Table.2 shows also that our non-Euclidean CNN-LSTM based model achieves competitive result to the state of the art (Shahrourdy et al., 2016) in terms of cross-subject accuracy. In fact, our model reaches 61.45% accuracy versus 62.93% in (Shahrourdy et al., 2016). For cross-view accuracy, our method outperforms the state of the art with 0.76% increase.

6 CONCLUSIONS

In this paper, we have proposed, for action recognition, to map skeleton sequences from the Riemannian manifold to linear spaces, previous to feature extraction and learning layers. We proposed a first non-Euclidean architecture based on CNNs to extract a compact representation of each skeletons frame. We then propose a second non-Euclidean temporally-aware architecture based on CNN-LSTM networks. We have tested the proposed approaches using two datasets, namely Parkinson's Vision-Based Pose Estimation dataset and NTU RGB+D dataset. Experimental results have shown the effectiveness of the proposed architectures compared to state of the art models. However, for future work, we are working 1) on integrating our method with state of the art architectures to consolidate its performance and 2) on improving the geometry awareness of deep learning architecture for action recognition by modifying the inner operations of the CNN network.

ACKNOWLEDGEMENTS

This work has been jointly supported by Talan Innovation Factory, Talan Tunisia, Talan Group. Talan is a French digital transformation Consulting Group, based in Paris, with offices in London, Geneva, Madrid, Luxembourg, New York, Chicago, Montreal, Toronto, Tunis, Rabat and Singapore. Talan Innovation Factory provides expertise relative to disruptive technologies such as Blockchain, Artificial Intelligence, Data Science and Internet of Things. In the frame of an academic-industry collaboration, Talan has been persistently contributing to this work by providing Hardware resources (Deep learning platform), mentoring and financial support.

REFERENCES

- Baccouche, M., Mamalet, F., Wolf, C., Garcia, C., and Baskurt, A. (2011). Sequential deep learning for human action recognition. In *Human Behavior Understanding - Second International Workshop, HBU 2011, Amsterdam, The Netherlands, November 16, 2011. Proceedings*, pages 29–39.
- Bronstein, M. M., Bruna, J., LeCun, Y., Szlam, A., and Vandergheynst, P. (2017). Geometric deep learning: Going beyond euclidean data. *IEEE Signal Process. Mag.*, 34(4):18–42.
- Ciresan, D. C., Meier, U., and Schmidhuber, J. (2012). Multi-column deep neural networks for image classification. In *2012 IEEE Conference on Computer Vision and Pattern Recognition, Providence, RI, USA, June 16-21, 2012*, pages 3642–3649.
- Cohen, T. S., Geiger, M., Koehler, J., and Welling, M. (2018). Spherical cnns.
- Donahue, J., Hendricks, L. A., Rohrbach, M., Venugopalan, S., Guadarrama, S., Saenko, K., and Darrell, T. (2017a). Long-term recurrent convolutional networks for visual recognition and description. *IEEE Trans. Pattern Anal. Mach. Intell.*, 39(4):677–691.
- Donahue, J., Hendricks, L. A., Rohrbach, M., Venugopalan, S., Guadarrama, S., Saenko, K., and Darrell, T. (2017b). Long-term recurrent convolutional networks for visual recognition and description. *IEEE Trans. Pattern Anal. Mach. Intell.*, 39(4):677–691.
- Du, Y., Wang, W., and Wang, L. (2015). Hierarchical recurrent neural network for skeleton based action recognition. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2015, Boston, MA, USA, June 7-12, 2015*, pages 1110–1118.
- Graves, A. (2012). *Supervised Sequence Labelling with Recurrent Neural Networks*, volume 385 of *Studies in Computational Intelligence*. Springer.
- Graves, A., Mohamed, A., and Hinton, G. E. (2013). Speech recognition with deep recurrent neural networks. In *IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP 2013, Vancouver, BC, Canada, May 26-31, 2013*, pages 6645–6649.
- Ke, Q., An, S., Bennamoun, M., Sohel, F. A., and Boussaïd, F. (2017). Skeletonnet: Mining deep part features for 3-d action recognition. *IEEE Signal Process. Lett.*, 24(6):731–735.
- Ke, Q. and Li, Y. (2014). Is rotation a nuisance in shape recognition? In *2014 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2014, Columbus, OH, USA, June 23-28, 2014*, pages 4146–4153.
- Krizhevsky, A., Sutskever, I., and Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems 25: 26th Annual Conference on Neural Information Processing Systems 2012. Proceedings of a meeting held December 3-6, 2012, Lake Tahoe, Nevada, United States*, pages 1106–1114.
- LeCun, Y. and Bengio, Y. (1998). The handbook of brain theory and neural networks. chapter Convolutional

- Networks for Images, Speech, and Time Series, pages 255–258. MIT Press, Cambridge, MA, USA.
- Lefebvre, G., Berlemont, S., Mamalet, F., and Garcia, C. (2013). BLSTM-RNN based 3d gesture classification. In *Artificial Neural Networks and Machine Learning - ICANN 2013 - 23rd International Conference on Artificial Neural Networks, Sofia, Bulgaria, September 10-13, 2013. Proceedings*, pages 381–388.
- Li, M. H., Mestre, T. A., Fox, S. H., and Taati, B. (2017). Vision-based assessment of parkinsonism and levodopa-induced dyskinesia with deep learning pose estimation. *CoRR*, abs/1707.09416.
- Liu, J., Shahroudy, A., Xu, D., and Wang, G. (2016). Spatio-temporal LSTM with trust gates for 3d human action recognition. In *Computer Vision - ECCV 2016 - 14th European Conference, Amsterdam, The Netherlands, October 11-14, 2016, Proceedings, Part III*, pages 816–833.
- Masci, J., Boscaini, D., Bronstein, M. M., and Vandergheynst, P. (2015). Geodesic convolutional neural networks on riemannian manifolds. In *2015 IEEE International Conference on Computer Vision Workshop, ICCV Workshops 2015, Santiago, Chile, December 7-13, 2015*, pages 832–840.
- Ng, J. Y., Hausknecht, M. J., Vijayanarasimhan, S., Vinyals, O., Monga, R., and Toderici, G. (2015). Beyond short snippets: Deep networks for video classification. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2015, Boston, MA, USA, June 7-12, 2015*, pages 4694–4702.
- Nussbaum, R. L. and Ellis, C. E. (2003). Alzheimer’s disease and parkinson’s disease. *New england journal of medicine*, 348(14):1356–1364.
- Shahroudy, A., Liu, J., Ng, T., and Wang, G. (2016). NTU RGB+D: A large scale dataset for 3d human activity analysis. In *2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016, Las Vegas, NV, USA, June 27-30, 2016*, pages 1010–1019.
- Srivastava, N., Mansimov, E., and Salakhutdinov, R. (2015). Unsupervised learning of video representations using lstms. In *Proceedings of the 32nd International Conference on Machine Learning, ICML 2015, Lille, France, 6-11 July 2015*, pages 843–852.
- Veeriah, V., Zhuang, N., and Qi, G. (2015). Differential recurrent neural networks for action recognition. In *2015 IEEE International Conference on Computer Vision, ICCV 2015, Santiago, Chile, December 7-13, 2015*, pages 4041–4049.
- Vemulapalli, R., Arrate, F., and Chellappa, R. (2014). Human action recognition by representing 3d skeletons as points in a lie group. In *2014 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2014, Columbus, OH, USA, June 23-28, 2014*, pages 588–595.
- Wang, L., Xiong, Y., Wang, Z., Qiao, Y., Lin, D., Tang, X., and Gool, L. V. (2016). Temporal segment networks: Towards good practices for deep action recognition. In *Computer Vision - ECCV 2016 - 14th European Conference, Amsterdam, The Netherlands, October 11-14, 2016, Proceedings, Part VIII*, pages 20–36.
- Wei, S., Ramakrishna, V., Kanade, T., and Sheikh, Y. (2016). Convolutional pose machines. In *2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016, Las Vegas, NV, USA, June 27-30, 2016*, pages 4724–4732.
- Xiong, Y., Zhu, K., Lin, D., and Tang, X. (2015). Recognize complex events from static images by fusing deep channels. In *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2015, Boston, MA, USA, June 7-12, 2015*, pages 1600–1609.
- Zhu, W., Lan, C., Xing, J., Zeng, W., Li, Y., Shen, L., and Xie, X. (2016). Co-occurrence feature learning for skeleton based action recognition using regularized deep LSTM networks. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence, February 12-17, 2016, Phoenix, Arizona, USA*, pages 3697–3704.