

Get started

Open in app



Follow

606K Followers

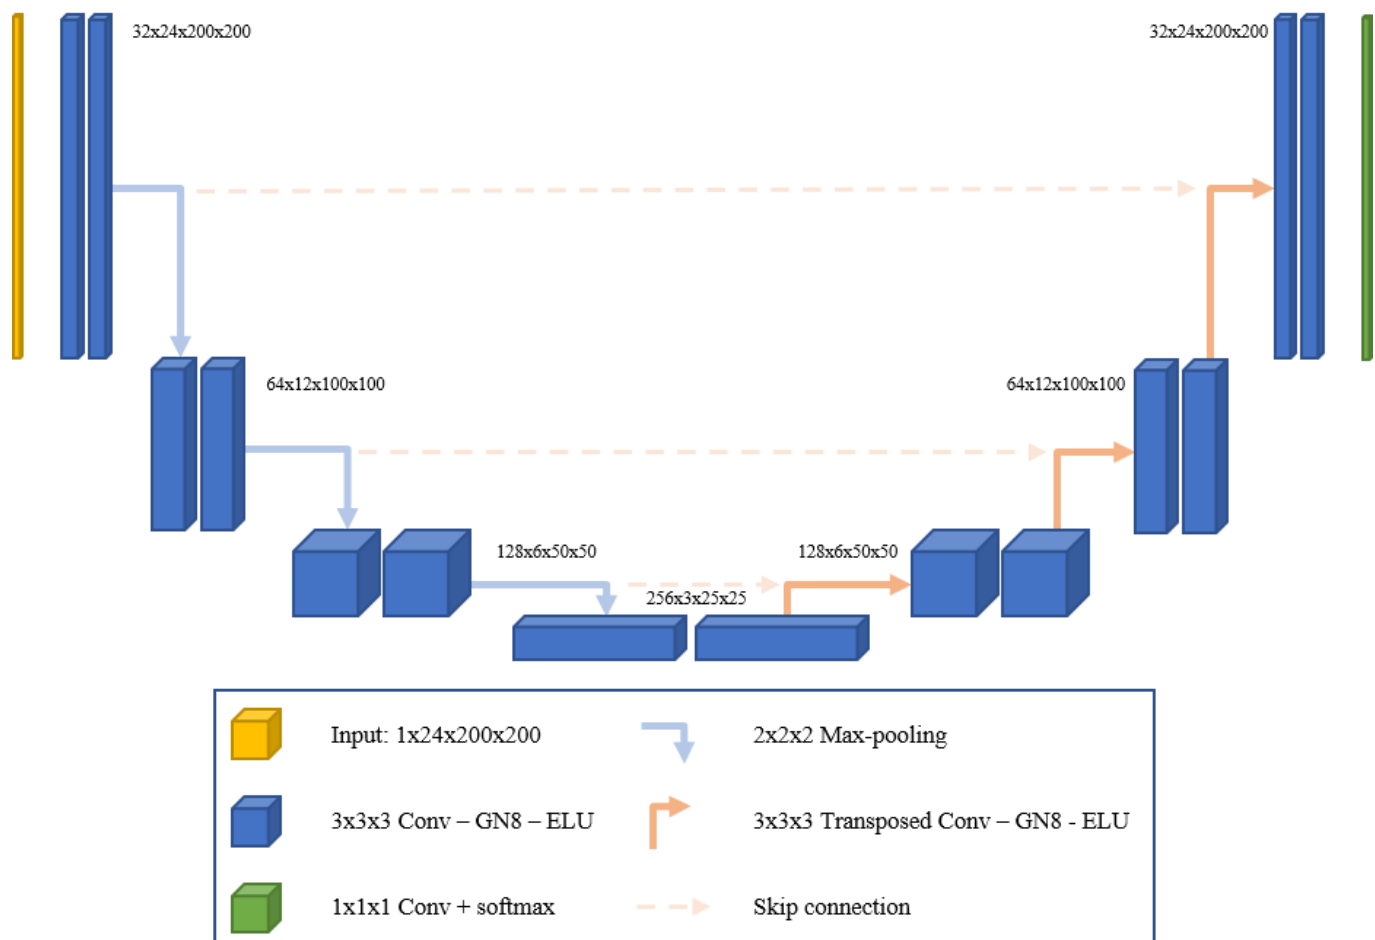


Creating and training a U-Net model with PyTorch for 2D & 3D semantic segmentation: Model building [2/4]

A guide to semantic segmentation with PyTorch and the U-Net



Johannes Schmidt Dec 2, 2020 · 6 min read



[Get started](#)[Open in app](#)

In the [previous chapter](#), we built a dataloader that picks up our images and performs some transformations and augmentations so that they can be fed in batches to a neural network like the U-Net. In this part, we focus on building a U-Net from scratch with the PyTorch library. The goal is to implement the U-Net in such a way, that important model configurations such as the activation function or the depth can be passed as arguments when creating the model.

About the U-Net

The U-Net is a convolutional neural network architecture that is designed for fast and precise segmentation of images. It has performed extremely well in several challenges and to this day, it is one of the most popular end-to-end architectures in the field of semantic segmentation.

We can split the network into two parts: The encoder path (backbone) and the decoder path. The encoder captures features at different scales of the images by using a traditional stack of convolutional and max pooling layers. Concretely speaking, a block in the encoder consists of the repeated use of two convolutional layers ($k=3, s=1$), each followed by a non-linearity layer, and a max-pooling layer ($k=2, s=2$). For every convolution block and its associated max pooling operation, the number of feature maps is doubled to ensure that the network can learn the complex structures effectively.

The decoder path is a symmetric expanding counterpart that uses transposed convolutions. This type of convolutional layer is an up-sampling method with trainable parameters and performs the reverse of (down)pooling layers such as the max pool. Similar to the encoder, each convolution block is followed by such an up-convolutional layer. The number of feature maps is halved in every block. Because recreating a segmentation mask from a small feature map is a rather difficult task for the network, the output after every up-convolutional layer is appended by the feature maps of the corresponding encoder block. The feature maps of the encoder layer are cropped if the dimensions exceed the one of the corresponding decoder layers.

In the end, the output passes another convolution layer ($k=1, s=1$) with the number of feature maps being equal to the number of defined labels. The result is a u-shaped

Get started

Open in app



The code

This code is based on

<https://github.com/ELEKTRONN/elektronn3/blob/master/elektronn3/models/unet.py>

© 2017 Martin Drawitsch, released under MIT License, which implements a configurable (2D/3D) U-Net with user-defined network depth and a few other improvements of the original architecture. They themselves actually used the 2D code from Jackson Huang <https://github.com/jaxony/unet-pytorch>.

Here is a simplified version of the code — saved in a file `unet.py`:

```

1  from torch import nn
2  import torch
3
4
5  @torch.jit.script
6  def autocrop(encoder_layer: torch.Tensor, decoder_layer: torch.Tensor):
7      """
8      Center-crops the encoder_layer to the size of the decoder_layer,
9      so that merging (concatenation) between levels/blocks is possible.
10     This is only necessary for input sizes != 2**n for 'same' padding and always required for '
11     """
12     if encoder_layer.shape[2:] != decoder_layer.shape[2:]:
13         ds = encoder_layer.shape[2:]
14         es = decoder_layer.shape[2:]
15         assert ds[0] >= es[0]
16         assert ds[1] >= es[1]
17         if encoder_layer.dim() == 4: # 2D
18             encoder_layer = encoder_layer[
19                 :,
20                 :,
21                 ((ds[0] - es[0]) // 2):((ds[0] + es[0]) // 2),
22                 ((ds[1] - es[1]) // 2):((ds[1] + es[1]) // 2)
23             ]
24         elif encoder_layer.dim() == 5: # 3D
25             assert ds[2] >= es[2]
26             encoder_layer = encoder_layer[
27                 :,
28                 :,
29                 ((ds[0] - es[0]) // 2):((ds[0] + es[0]) // 2),

```

Get started

Open in app



```
32         ]
33     return encoder_layer, decoder_layer
34
35
36 def conv_layer(dim: int):
37     if dim == 3:
38         return nn.Conv3d
39     elif dim == 2:
40         return nn.Conv2d
41
42
43 def get_conv_layer(in_channels: int,
44                   out_channels: int,
45                   kernel_size: int = 3,
46                   stride: int = 1,
47                   padding: int = 1,
48                   bias: bool = True,
49                   dim: int = 2):
50     return conv_layer(dim)(in_channels, out_channels, kernel_size=kernel_size, stride=stride, padding=padding, bias=bias)
51
52
53
54 def conv_transpose_layer(dim: int):
55     if dim == 3:
56         return nn.ConvTranspose3d
57     elif dim == 2:
58         return nn.ConvTranspose2d
59
60
61 def get_up_layer(in_channels: int,
62                out_channels: int,
63                kernel_size: int = 2,
64                stride: int = 2,
65                dim: int = 3,
66                up_mode: str = 'transposed',
67                ):
68     if up_mode == 'transposed':
69         return conv_transpose_layer(dim)(in_channels, out_channels, kernel_size=kernel_size, stride=stride, padding=padding, bias=bias)
70     else:
71         return nn.Upsample(scale_factor=2.0, mode=up_mode)
72
73
```

Get started

Open in app



```
77     elif dim == 2:
78         return nn.MaxPool2d
79
80
81     def get_maxpool_layer(kernel_size: int = 2,
82                           stride: int = 2,
83                           padding: int = 0,
84                           dim: int = 2):
85         return maxpool_layer(dim=dim)(kernel_size=kernel_size, stride=stride, padding=padding)
86
87
88     def get_activation(activation: str):
89         if activation == 'relu':
90             return nn.ReLU()
91         elif activation == 'leaky':
92             return nn.LeakyReLU(negative_slope=0.1)
93         elif activation == 'elu':
94             return nn.ELU()
95
96
97     def get_normalization(normalization: str,
98                           num_channels: int,
99                           dim: int):
100         if normalization == 'batch':
101             if dim == 3:
102                 return nn.BatchNorm3d(num_channels)
103             elif dim == 2:
104                 return nn.BatchNorm2d(num_channels)
105         elif normalization == 'instance':
106             if dim == 3:
107                 return nn.InstanceNorm3d(num_channels)
108             elif dim == 2:
109                 return nn.InstanceNorm2d(num_channels)
110         elif 'group' in normalization:
111             num_groups = int(normalization.partition('group')[-1]) # get the group size from string
112             return nn.GroupNorm(num_groups=num_groups, num_channels=num_channels)
113
114
115     class Concatenate(nn.Module):
116         def __init__(self):
117             super(Concatenate, self).__init__()
```

Get started

Open in app



```
121
122     return x
123
124
125 class DownBlock(nn.Module):
126     """
127     A helper Module that performs 2 Convolutions and 1 MaxPool.
128     An activation follows each convolution.
129     A normalization layer follows each convolution.
130     """
131
132     def __init__(self,
133                 in_channels: int,
134                 out_channels: int,
135                 pooling: bool = True,
136                 activation: str = 'relu',
137                 normalization: str = None,
138                 dim: str = 2,
139                 conv_mode: str = 'same'):
140         super().__init__()
141
142         self.in_channels = in_channels
143         self.out_channels = out_channels
144         self.pooling = pooling
145         self.normalization = normalization
146         if conv_mode == 'same':
147             self.padding = 1
148         elif conv_mode == 'valid':
149             self.padding = 0
150         self.dim = dim
151         self.activation = activation
152
153         # conv layers
154         self.conv1 = get_conv_layer(self.in_channels, self.out_channels, kernel_size=3, stride=
155                                     bias=True, dim=self.dim)
156         self.conv2 = get_conv_layer(self.out_channels, self.out_channels, kernel_size=3, stride=
157                                     bias=True, dim=self.dim)
158
159         # pooling layer
160         if self.pooling:
161             self.pool = get_maxpool_layer(kernel_size=2, stride=2, padding=0, dim=self.dim)
162
```

Get started

Open in app



```

166
167     # normalization layers
168     if self.normalization:
169         self.norm1 = get_normalization(normalization=self.normalization, num_channels=self.c
170                                         dim=self.dim)
171         self.norm2 = get_normalization(normalization=self.normalization, num_channels=self.c
172                                         dim=self.dim)
173
174     def forward(self, x):
175         y = self.conv1(x) # convolution 1
176         y = self.act1(y) # activation 1
177         if self.normalization:
178             y = self.norm1(y) # normalization 1
179         y = self.conv2(y) # convolution 2
180         y = self.act2(y) # activation 2
181         if self.normalization:
182             y = self.norm2(y) # normalization 2
183
184         before_pooling = y # save the outputs before the pooling operation
185         if self.pooling:
186             y = self.pool(y) # pooling
187         return y, before_pooling
188
189
190 class UpBlock(nn.Module):
191     """
192     A helper Module that performs 2 Convolutions and 1 UpConvolution/Upsample.
193     An activation follows each convolution.
194     A normalization layer follows each convolution.
195     """
196
197     def __init__(self,
198                 in_channels: int,
199                 out_channels: int,
200                 activation: str = 'relu',
201                 normalization: str = None,
202                 dim: int = 3,
203                 conv_mode: str = 'same',
204                 up_mode: str = 'transposed'
205                 ):
206         super().__init__()
207

```

Get started

Open in app



```

210     self.normalization = normalization
211     if conv_mode == 'same':
212         self.padding = 1
213     elif conv_mode == 'valid':
214         self.padding = 0
215     self.dim = dim
216     self.activation = activation
217     self.up_mode = up_mode
218
219     # upconvolution/upsample layer
220     self.up = get_up_layer(self.in_channels, self.out_channels, kernel_size=2, stride=2, dim=
221                           up_mode=self.up_mode)
222
223     # conv layers
224     self.conv0 = get_conv_layer(self.in_channels, self.out_channels, kernel_size=1, stride=
225                                bias=True, dim=self.dim)
226     self.conv1 = get_conv_layer(2 * self.out_channels, self.out_channels, kernel_size=3, st
227                                padding=self.padding,
228                                bias=True, dim=self.dim)
229     self.conv2 = get_conv_layer(self.out_channels, self.out_channels, kernel_size=3, stride
230                                bias=True, dim=self.dim)
231
232     # activation layers
233     self.act0 = get_activation(self.activation)
234     self.act1 = get_activation(self.activation)
235     self.act2 = get_activation(self.activation)
236
237     # normalization layers
238     if self.normalization:
239         self.norm0 = get_normalization(normalization=self.normalization, num_channels=self.o
240                                       dim=self.dim)
241         self.norm1 = get_normalization(normalization=self.normalization, num_channels=self.o
242                                       dim=self.dim)
243         self.norm2 = get_normalization(normalization=self.normalization, num_channels=self.o
244                                       dim=self.dim)
245
246     # concatenate layer
247     self.concat = Concatenate()
248
249     def forward(self, encoder_layer, decoder_layer):
250         """ Forward pass
251         Arguments:

```


[Get started](#)[Open in app](#)

```
255         up_layer = self.up(decoder_layer) # up-convolution/up-sampling
256         cropped_encoder_layer, dec_layer = autocrop(encoder_layer, up_layer) # cropping
257
258         if self.up_mode != 'transposed':
259             # We need to reduce the channel dimension with a conv layer
260             up_layer = self.conv0(up_layer) # convolution 0
261         up_layer = self.act0(up_layer) # activation 0
262         if self.normalization:
263             up_layer = self.norm0(up_layer) # normalization 0
264
265         merged_layer = self.concat(up_layer, cropped_encoder_layer) # concatenation
266         y = self.conv1(merged_layer) # convolution 1
267         y = self.act1(y) # activation 1
268         if self.normalization:
269             y = self.norm1(y) # normalization 1
270         y = self.conv2(y) # convolution 2
271         y = self.act2(y) # activation 2
272         if self.normalization:
273             y = self.norm2(y) # normalization 2
274         return y
275
276
277     class UNet(nn.Module):
278         def __init__(self,
279                     in_channels: int = 1,
280                     out_channels: int = 2,
281                     n_blocks: int = 4,
282                     start_filters: int = 32,
283                     activation: str = 'relu',
284                     normalization: str = 'batch',
285                     conv_mode: str = 'same',
286                     dim: int = 2,
287                     up_mode: str = 'transposed'
288                 ):
289             super().__init__()
290
291             self.in_channels = in_channels
292             self.out_channels = out_channels
293             self.n_blocks = n_blocks
294             self.start_filters = start_filters
295             self.activation = activation
296             self.normalization = normalization
```

Get started

Open in app



```
299         self.up_mode = up_mode
300
301     self.down_blocks = []
302     self.up_blocks = []
303
304     # create encoder path
305     for i in range(self.n_blocks):
306         num_filters_in = self.in_channels if i == 0 else num_filters_out
307         num_filters_out = self.start_filters * (2 ** i)
308         pooling = True if i < self.n_blocks - 1 else False
309
310         down_block = DownBlock(in_channels=num_filters_in,
311                                out_channels=num_filters_out,
312                                pooling=pooling,
313                                activation=self.activation,
314                                normalization=self.normalization,
315                                conv_mode=self.conv_mode,
316                                dim=self.dim)
317
318         self.down_blocks.append(down_block)
319
320     # create decoder path (requires only n_blocks-1 blocks)
321     for i in range(n_blocks - 1):
322         num_filters_in = num_filters_out
323         num_filters_out = num_filters_in // 2
324
325         up_block = UpBlock(in_channels=num_filters_in,
326                            out_channels=num_filters_out,
327                            activation=self.activation,
328                            normalization=self.normalization,
329                            conv_mode=self.conv_mode,
330                            dim=self.dim,
331                            up_mode=self.up_mode)
332
333         self.up_blocks.append(up_block)
334
335     # final convolution
336     self.conv_final = get_conv_layer(num_filters_out, self.out_channels, kernel_size=1, str:
337                                       bias=True, dim=self.dim)
338
339     # add the list of modules to current module
340     self.down_blocks = nn.ModuleList(self.down_blocks)
```

Get started

Open in app



```

344         self.initialize_parameters()
345
346     @staticmethod
347     def weight_init(module, method, **kwargs):
348         if isinstance(module, (nn.Conv3d, nn.Conv2d, nn.ConvTranspose3d, nn.ConvTranspose2d)):
349             method(module.weight, **kwargs) # weights
350
351     @staticmethod
352     def bias_init(module, method, **kwargs):
353         if isinstance(module, (nn.Conv3d, nn.Conv2d, nn.ConvTranspose3d, nn.ConvTranspose2d)):
354             method(module.bias, **kwargs) # bias
355
356     def initialize_parameters(self,
357                             method_weights=nn.init.xavier_uniform_,
358                             method_bias=nn.init.zeros_,
359                             kwargs_weights={},
360                             kwargs_bias={}
361                             ):
362         for module in self.modules():
363             self.weight_init(module, method_weights, **kwargs_weights) # initialize weights
364             self.bias_init(module, method_bias, **kwargs_bias) # initialize bias
365
366     def forward(self, x: torch.tensor):
367         encoder_output = []
368
369         # Encoder pathway
370         for module in self.down_blocks:
371             x, before_pooling = module(x)
372             encoder_output.append(before_pooling)
373
374         # Decoder pathway
375         for i, module in enumerate(self.up_blocks):
376             before_pool = encoder_output[-(i + 2)]
377             x = module(before_pool, x)
378
379         x = self.conv_final(x)
380
381         return x
382
383     def __repr__(self):
384         attributes = {attr_key: self.__dict__[attr_key] for attr_key in self.__dict__.keys() if
385         d = {self.__class__.__name__: attributes}

```


[Get started](#)[Open in app](#)

UNET.py hosted with ❤ by GitHub

[view raw](#)


I will not go into detail here, but rather just mention important design choices. It can be useful to view the architecture in repeating blocks in the encoder but also in the decoder path. As you can see in `UNET.py` the `DownBlock` and the `UpBlock` help to build the architecture. Both use smaller helper functions that return the correct layer, depending on what arguments are passed, e.g. if a 2D (`dim=2`) or 3D (`dim=3`) network is wanted. The number of blocks is defined by the `depth` of the network.

A `DownBlock` generally has the following scheme:



DownBlock — Image by Johannes Schmidt

A `UpBlock` has the following layers:



UpBlock — Image by Johannes Schmidt

For our `Unet` class we just need to combine these blocks and make sure that the correct layers from the encoder are concatenated to the decoder (skip pathways). These layers have to be cropped if their sizes do not match with the corresponding layers from the decoder. In such cases, the `autocrop` function is used. For merging, I concatenate along the channel dimension (see `Concatenate`). Instead of transposed convolutions we could also use upsampling layers (interpolation methods) that are followed by a 1x1 or 3x3 convolution block to reduce the channel dimension. Using interpolation generally gets rid of the checkerboard artifact. For 3D input consider using trilinear interpolation.

At the end we just need to think about the parameter initialization. By default, the weights are initialized with `torch.nn.init.xavier_uniform_` and the biases are initialized with zeros using `torch.nn.init.zeros_`.

[Get started](#)[Open in app](#)

functions.

Creating a U-Net model

Let's create such a model and use it to make a prediction on some random input:

```
from unet import UNet

model = UNet(in_channels=1,
             out_channels=2,
             n_blocks=4,
             start_filters=32,
             activation='relu',
             normalization='batch',
             conv_mode='same',
             dim=2)

x = torch.randn(size=(1, 1, 512, 512), dtype=torch.float32)
with torch.no_grad():
    out = model(x)

print(f'Out: {out.shape}')
```

This will give us:

```
Out: torch.Size([1, 2, 512, 512])
```

To check whether our model is correct, we can get the model's summary with this package [pytorch-summary](#):

```
from torchsummary import summary
summary = summary(model, (1, 512, 512))
```

which prints out a summary like this:

```
1  =====
2  Layer (type:depth-idx)           Output Shape           Param #
```

Get started

Open in app




6			└─Conv2d: 3-1	[-1, 32, 512, 512]	320
7			└─ReLU: 3-2	[-1, 32, 512, 512]	--
8			└─BatchNorm2d: 3-3	[-1, 32, 512, 512]	64
9			└─Conv2d: 3-4	[-1, 32, 512, 512]	9,248
10			└─ReLU: 3-5	[-1, 32, 512, 512]	--
11			└─BatchNorm2d: 3-6	[-1, 32, 512, 512]	64
12			└─MaxPool2d: 3-7	[-1, 32, 256, 256]	--
13			└─DownBlock: 2-2	[-1, 64, 128, 128]	--
14			└─Conv2d: 3-8	[-1, 64, 256, 256]	18,496
15			└─ReLU: 3-9	[-1, 64, 256, 256]	--
16			└─BatchNorm2d: 3-10	[-1, 64, 256, 256]	128
17			└─Conv2d: 3-11	[-1, 64, 256, 256]	36,928
18			└─ReLU: 3-12	[-1, 64, 256, 256]	--
19			└─BatchNorm2d: 3-13	[-1, 64, 256, 256]	128
20			└─MaxPool2d: 3-14	[-1, 64, 128, 128]	--
21			└─DownBlock: 2-3	[-1, 128, 64, 64]	--
22			└─Conv2d: 3-15	[-1, 128, 128, 128]	73,856
23			└─ReLU: 3-16	[-1, 128, 128, 128]	--
24			└─BatchNorm2d: 3-17	[-1, 128, 128, 128]	256
25			└─Conv2d: 3-18	[-1, 128, 128, 128]	147,584
26			└─ReLU: 3-19	[-1, 128, 128, 128]	--
27			└─BatchNorm2d: 3-20	[-1, 128, 128, 128]	256
28			└─MaxPool2d: 3-21	[-1, 128, 64, 64]	--
29			└─DownBlock: 2-4	[-1, 256, 64, 64]	--
30			└─Conv2d: 3-22	[-1, 256, 64, 64]	295,168
31			└─ReLU: 3-23	[-1, 256, 64, 64]	--
32			└─BatchNorm2d: 3-24	[-1, 256, 64, 64]	512
33			└─Conv2d: 3-25	[-1, 256, 64, 64]	590,080
34			└─ReLU: 3-26	[-1, 256, 64, 64]	--
35			└─BatchNorm2d: 3-27	[-1, 256, 64, 64]	512
36			└─ModuleList: 1	[]	--
37			└─UpBlock: 2-5	[-1, 128, 128, 128]	--
38			└─ConvTranspose2d: 3-28	[-1, 128, 128, 128]	131,200
39			└─ReLU: 3-29	[-1, 128, 128, 128]	--
40			└─BatchNorm2d: 3-30	[-1, 128, 128, 128]	256
41			└─Concatenate: 3-31	[-1, 256, 128, 128]	--
42			└─Conv2d: 3-32	[-1, 128, 128, 128]	295,040
43			└─ReLU: 3-33	[-1, 128, 128, 128]	--
44			└─BatchNorm2d: 3-34	[-1, 128, 128, 128]	256
45			└─Conv2d: 3-35	[-1, 128, 128, 128]	147,584
46			└─ReLU: 3-36	[-1, 128, 128, 128]	--

Get started

Open in app

50			└ReLU: 3-39	[-1, 64, 256, 256]	--
51			└BatchNorm2d: 3-40	[-1, 64, 256, 256]	128
52			└Concatenate: 3-41	[-1, 128, 256, 256]	--
53			└Conv2d: 3-42	[-1, 64, 256, 256]	73,792
54			└ReLU: 3-43	[-1, 64, 256, 256]	--
55			└BatchNorm2d: 3-44	[-1, 64, 256, 256]	128
56			└Conv2d: 3-45	[-1, 64, 256, 256]	36,928
57			└ReLU: 3-46	[-1, 64, 256, 256]	--
58			└BatchNorm2d: 3-47	[-1, 64, 256, 256]	128
59			└UpBlock: 2-7	[-1, 32, 512, 512]	--
60			└ConvTranspose2d: 3-48	[-1, 32, 512, 512]	8,224
61			└ReLU: 3-49	[-1, 32, 512, 512]	--
62			└BatchNorm2d: 3-50	[-1, 32, 512, 512]	64
63			└Concatenate: 3-51	[-1, 64, 512, 512]	--
64			└Conv2d: 3-52	[-1, 32, 512, 512]	18,464
65			└ReLU: 3-53	[-1, 32, 512, 512]	--
66			└BatchNorm2d: 3-54	[-1, 32, 512, 512]	64
67			└Conv2d: 3-55	[-1, 32, 512, 512]	9,248
68			└ReLU: 3-56	[-1, 32, 512, 512]	--
69			└BatchNorm2d: 3-57	[-1, 32, 512, 512]	64
70			└Conv2d: 1-1	[-1, 2, 512, 512]	66
71	=====				
72	Total params: 1,928,322				
73	Trainable params: 1,928,322				
74	Non-trainable params: 0				
75	Total mult-adds (M): -14267.05				
76	=====				
77	Input size (MB): 1.00				
78	Forward/backward pass size (MB): 1156.00				
79	Params size (MB): 7.36				
80	Estimated Total Size (MB): 1164.36				
81	=====				

unet_model_summary.txt hosted with  by GitHub

view raw

About input sizes

To ensure correct semantic concatenations, it is advised to use input sizes that return even spatial dimensions in every block but the last in the encoder. For example: An input size of 120^2 gives intermediate output shapes of $[60^2, 30^2, 15^2]$ in the encoder path for a U-Net with `depth=4` . A U-Net with `depth=5` with the same input size is not

[Get started](#)[Open in app](#)

To make our lives easier, we can numerically compute the maximum network depth for a given input dimension with a simple function:

```
shape = 1920

def compute_max_depth(shape, max_depth=10, print_out=True):
    shapes = []
    shapes.append(shape)
    for level in range(1, max_depth):
        if shape % 2 ** level == 0 and shape / 2 ** level > 1:
            shapes.append(shape / 2 ** level)
            if print_out:
                print(f'Level {level}: {shape / 2 ** level}')
        else:
            if print_out:
                print(f'Max-level: {level - 1}')
            break

    return shapes

out = compute_max_depth(shape, print_out=True, max_depth=10)
```

This will output

```
Level 1: 960.0
Level 2: 480.0
Level 3: 240.0
Level 4: 120.0
Level 5: 60.0
Level 6: 30.0
Level 7: 15.0
Max-level: 7
```

which tells us that that we can design a U-Net as deep as this without having to worry about semantic mismatches. Conversely, we can also numerically determine the possible input shapes dimensions for a given depth:

[Get started](#)[Open in app](#)

```
def compute_possible_shapes(low, high, depth):
    possible_shapes = {}
    for shape in range(low, high + 1):
        shapes = compute_max_depth(shape,
                                    max_depth=depth,
                                    print_out=False)
        if len(shapes) == depth:
            possible_shapes[shape] = shapes

    return possible_shapes

possible_shapes = compute_possible_shapes(low, high, depth)
```

This will output

```
{256: [256, 128.0, 64.0, 32.0, 16.0, 8.0, 4.0, 2.0],
 384: [384, 192.0, 96.0, 48.0, 24.0, 12.0, 6.0, 3.0],
 512: [512, 256.0, 128.0, 64.0, 32.0, 16.0, 8.0, 4.0]}
```

which tells us that we can have 3 different input shapes with such a level 8 U-Net architecture. But I dare to say that such a network with this input size is probably not useful in practice.

Summary

In this part we created a configurable UNet model for the purpose of semantic segmentation. Now that we have built our model, it is time to create a training loop in the [next chapter](#).

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

You'll need to sign in or create an account to receive this

[Get started](#)[Open in app](#)[Unet](#)[Python](#)[Pytorch](#)[Deep Learning](#)[Semantic Segmentation](#)[About](#) [Help](#) [Legal](#)

Get the Medium app

