

[Get started](#)[Open in app](#)[Follow](#)

605K Followers



Creating and training a U-Net model with PyTorch for 2D & 3D semantic segmentation: Training [3/4]

A guide to semantic segmentation with PyTorch and the U-Net



Johannes Schmidt · Dec 5, 2020 · 7 min read

```
In [40]: from visual import plot_training
fig = plot_training(training_losses,
                    validation_losses,
                    learning_rates,
                    gaussian=True,
                    sigma=1,
                    figsize=(10, 4))
```

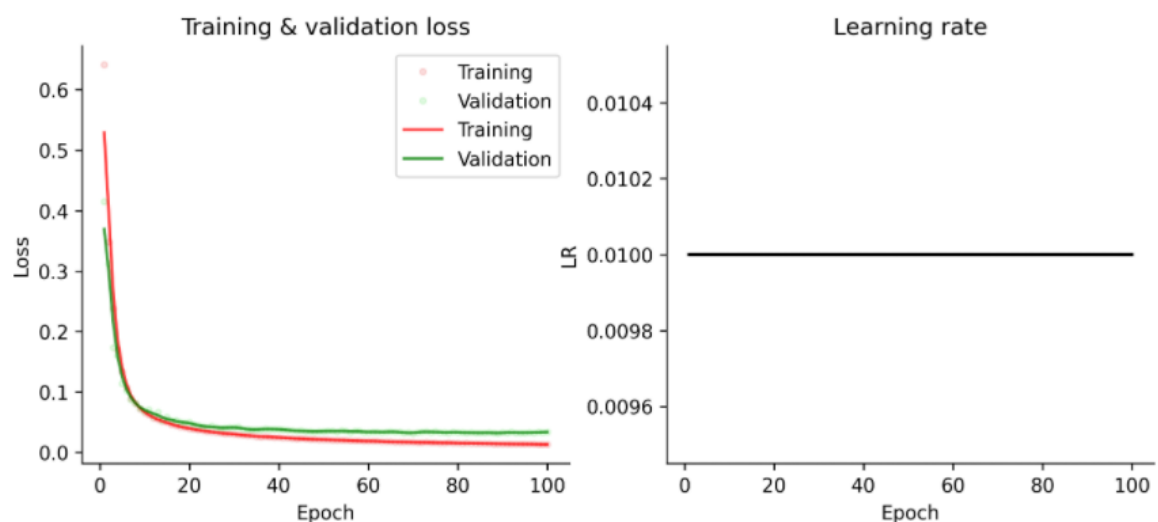


Image by author

[Get started](#)[Open in app](#)

`Trainer` class and save it in `trainer.py`. The Jupyter notebook can be found [here](#). The idea is that we can instantiate a `Trainer` object with parameters such as the `model`, a `criterion` etc. and then call it's class method `run_trainer()` to start training. This method will output the accumulated training loss, the validation loss, and the learning rate that was used for training. Here is the code:

```
1  import numpy as np
2  import torch
3
4
5  class Trainer:
6      def __init__(self,
7                  model: torch.nn.Module,
8                  device: torch.device,
9                  criterion: torch.nn.Module,
10                 optimizer: torch.optim.Optimizer,
11                 training_DataLoader: torch.utils.data.Dataset,
12                 validation_DataLoader: torch.utils.data.Dataset = None,
13                 lr_scheduler: torch.optim.lr_scheduler = None,
14                 epochs: int = 100,
15                 epoch: int = 0,
16                 notebook: bool = False
17             ):
18
19         self.model = model
20         self.criterion = criterion
21         self.optimizer = optimizer
22         self.lr_scheduler = lr_scheduler
23         self.training_DataLoader = training_DataLoader
24         self.validation_DataLoader = validation_DataLoader
25         self.device = device
26         self.epochs = epochs
27         self.epoch = epoch
28         self.notebook = notebook
29
30         self.training_loss = []
31         self.validation_loss = []
32         self.learning_rate = []
33
34     def run_trainer(self):
```

Get started

Open in app



```

37         from tqdm.notebook import tqdm, trange
38     else:
39         from tqdm import tqdm, trange
40
41     progressbar = trange(self.epochs, desc='Progress')
42     for i in progressbar:
43         """Epoch counter"""
44         self.epoch += 1 # epoch counter
45
46         """Training block"""
47         self._train()
48
49         """Validation block"""
50         if self.validation_DataLoader is not None:
51             self._validate()
52
53         """Learning rate scheduler block"""
54         if self.lr_scheduler is not None:
55             if self.validation_DataLoader is not None and self.lr_scheduler.__class__.__name__ == 'StepLR':
56                 self.lr_scheduler.batch(self.validation_loss[i]) # learning rate scheduler
57             else:
58                 self.lr_scheduler.batch() # learning rate scheduler step
59     return self.training_loss, self.validation_loss, self.learning_rate
60
61     def _train(self):
62
63         if self.notebook:
64             from tqdm.notebook import tqdm, trange
65         else:
66             from tqdm import tqdm, trange
67
68         self.model.train() # train mode
69         train_losses = [] # accumulate the losses here
70         batch_iter = tqdm(enumerate(self.training_DataLoader), 'Training', total=len(self.training_DataLoader),
71                             leave=False)
72
73         for i, (x, y) in batch_iter:
74             input, target = x.to(self.device), y.to(self.device) # send to device (GPU or CPU)
75             self.optimizer.zero_grad() # zerograd the parameters
76             out = self.model(input) # one forward pass
77             loss = self.criterion(out, target) # calculate loss
78             loss_value = loss.item()

```

Get started

Open in app



```
82
83         batch_iter.set_description(f'Training: (loss {loss_value:.4f})') # update progress
84
85     self.training_loss.append(np.mean(train_losses))
86     self.learning_rate.append(self.optimizer.param_groups[0]['lr'])
87
88     batch_iter.close()
89
90     def _validate(self):
91
92         if self.notebook:
93             from tqdm.notebook import tqdm, trange
94         else:
95             from tqdm import tqdm, trange
96
97         self.model.eval() # evaluation mode
98         valid_losses = [] # accumulate the losses here
99         batch_iter = tqdm(enumerate(self.validation_DataLoader), 'Validation', total=len(self.v
100                             leave=False)
101
102         for i, (x, y) in batch_iter:
103             input, target = x.to(self.device), y.to(self.device) # send to device (GPU or CPU)
104
105             with torch.no_grad():
106                 out = self.model(input)
107                 loss = self.criterion(out, target)
108                 loss_value = loss.item()
109                 valid_losses.append(loss_value)
110
111             batch_iter.set_description(f'Validation: (loss {loss_value:.4f})')
112
113         self.validation_loss.append(np.mean(valid_losses))
114
115         batch_iter.close()
```

trainer.py hosted with ❤ by GitHub

[view raw](#)

In order to create a trainer object the following parameters are required:

- `model` : e.g. the U-Net

[Get started](#)[Open in app](#)

- `criterion` : loss function (e.g. `CrossEntropyLoss`, `DiceCoefficientLoss`)
- `optimizer` : e.g. `SGD`
- `training_DataLoader` : a training dataloader
- `validation_DataLoader` : a validation dataloader
- `lr_scheduler` : a learning rate scheduler (optional)
- `epochs` : The number of epochs we want to train
- `epoch` : The epoch number from where training should start

Training can then be started with the class method `run_trainer()`. Since training is usually performed with a training and a validation phase, `_train()` and `_validate()` are two functions that are run once for every epoch we train with `run_trainer()` (line 33–53). If we have a `lr_scheduler`, we also perform a step with the `lr_scheduler`. To visualize the progress of training, I included a progress bar with the library `tqdm`. Now let's take a closer look on what happens when calling `_train()` and `_validate()`. If you are familiar with using PyTorch for network training, there is probably nothing new here.

In `_train()` we basically just iterate over our training dataloader and send our batches through the network in train mode (line 56–64). We then use this output together with our target to compute the loss with the loss function for the current batch (line 65). The computed loss is then appended in a temporary list (line 66–67). Based on the computed gradients, we perform a backward pass and a step with our optimizer to update the model's parameters (line 68–69). At the end we update our progress bar for the training phase to show the current loss (line 71). The function outputs the mean of the temporary loss list and the learning rate that was used.

In `_validate()`, similar to `_train()`, we iterate over our validation dataloader, send our batches through the network in validation mode and compute the loss. This time, without computing the gradients and without performing a backward pass (line 78–97).

[Get started](#)[Open in app](#)

Jupyter notebook.

```
1  # Imports
2  import pathlib
3  import torch
4
5  import albumentations
6  import numpy as np
7  from sklearn.model_selection import train_test_split
8  from torch.utils.data import DataLoader
9  from skimage.transform import resize
10 from customdatasets import SegmentationDataSet1
11 from transformations import ComposeDouble, AlbuSeg2d, FunctionWrapperDouble, normalize_01, crea
12 from unet import UNet
13 from trainer import Trainer
14
15
16 # root directory
17 root = pathlib.Path.cwd() / 'Carvana'
18
19
20 def get_filenames_of_path(path: pathlib.Path, ext: str = '*'):
21     """Returns a list of files in a directory/path. Uses pathlib."""
22     filenames = [file for file in path.glob(ext) if file.is_file()]
23     return filenames
24
25
26 # input and target files
27 inputs = get_filenames_of_path(root / 'Input')
28 targets = get_filenames_of_path(root / 'Target')
29
30 # training transformations and augmentations
31 transforms_training = ComposeDouble([
32     FunctionWrapperDouble(resize,
33                           input=True,
34                           target=False,
35                           output_shape=(128, 128, 3)),
36     FunctionWrapperDouble(resize,
37                           input=False,
38                           target=True,
39                           output_shape=(128, 128),
```

Get started

Open in app



```
43     AlbuSeg2d(albumentations.HorizontalFlip(p=0.5)),
44     FunctionWrapperDouble(create_dense_target, input=False, target=True),
45     FunctionWrapperDouble(np.moveaxis, input=True, target=False, source=-1, destination=0),
46     FunctionWrapperDouble(normalize_01)
47 ])
48
49 # validation transformations
50 transforms_validation = ComposeDouble([
51     FunctionWrapperDouble(resize,
52                           input=True,
53                           target=False,
54                           output_shape=(128, 128, 3)),
55     FunctionWrapperDouble(resize,
56                           input=False,
57                           target=True,
58                           output_shape=(128, 128),
59                           order=0,
60                           anti_aliasing=False,
61                           preserve_range=True),
62     FunctionWrapperDouble(create_dense_target, input=False, target=True),
63     FunctionWrapperDouble(np.moveaxis, input=True, target=False, source=-1, destination=0),
64     FunctionWrapperDouble(normalize_01)
65 ])
66
67 # random seed
68 random_seed = 42
69
70 # split dataset into training set and validation set
71 train_size = 0.8 # 80:20 split
72
73 inputs_train, inputs_valid = train_test_split(
74     inputs,
75     random_state=random_seed,
76     train_size=train_size,
77     shuffle=True)
78
79 targets_train, targets_valid = train_test_split(
80     targets,
81     random_state=random_seed,
82     train_size=train_size,
83     shuffle=True)
84
```

Get started

Open in app



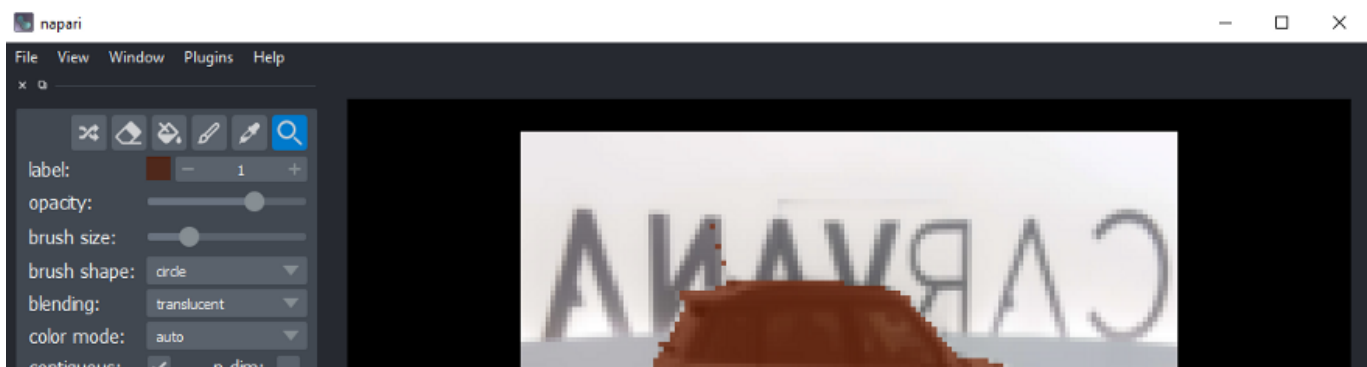
```
87
88 # dataset training
89 dataset_train = SegmentationDataSet1(inputs=inputs_train,
90                                     targets=targets_train,
91                                     transform=transforms_training)
92
93 # dataset validation
94 dataset_valid = SegmentationDataSet1(inputs=inputs_valid,
95                                     targets=targets_valid,
96                                     transform=transforms_validation)
97
98 # dataloader training
99 dataloader_training = DataLoader(dataset=dataset_train,
100                                 batch_size=2,
101                                 shuffle=True)
102
103 # dataloader validation
104 dataloader_validation = DataLoader(dataset=dataset_valid,
105                                   batch_size=2,
106                                   shuffle=True)
```

customdataset1.py hosted with ❤ by GitHub

view raw

Please note that I resize the images to 128x128x3 using `skimage.transform.resize()` to speed up training. This will generate batches of images that look like this:

```
from visual import DatasetViewer
dataset_viewer_training = DatasetViewer(dataset_train)
dataset_viewer_training.napari()
```



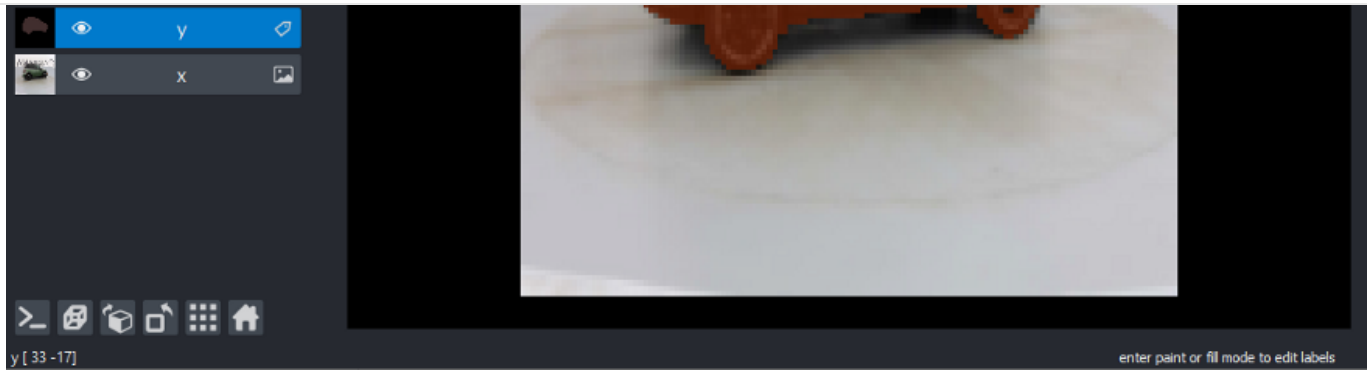
[Get started](#)[Open in app](#)

Image by author

I can then instantiate the `Trainer` object and start training:

```
1  # device
2  if torch.cuda.is_available():
3      device = torch.device('cuda')
4  else:
5      torch.device('cpu')
6
7  # model
8  model = UNet(in_channels=3,
9              out_channels=2,
10             n_blocks=4,
11             start_filters=32,
12             activation='relu',
13             normalization='batch',
14             conv_mode='same',
15             dim=2).to(device)
16
17 # criterion
18 criterion = torch.nn.CrossEntropyLoss()
19
20 # optimizer
21 optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
22
23 # trainer
24 trainer = Trainer(model=model,
25                  device=device,
26                  criterion=criterion,
27                  optimizer=optimizer,
28                  training_DataLoader=dataloader_training,
```

[Get started](#)[Open in app](#)

```
32         epoch=0,  
33         notebook=True)  
34  
35 # start training  
36 training_losses, validation_losses, lr_rates = trainer.run_trainer()
```

seg_start_training.py hosted with ❤ by GitHub

[view raw](#)

Training will look something like this:



```
In [*]: # start training  
training_losses, validation_losses, learning_rates = trainer.run_trainer()  
  
Progress: 50%  1/2 [00:39<00:39, 39.15s/it]  
  
Training: (loss 0.0770): 26%  10/38 [00:08<00:23, 1.21it/s]
```

Image by author

Improve the data generator

Although training was performed on a NVIDIA 1070, it took 1:19 min to train 2 epochs with only 96 images (size 128x128x3) for each epoch. Why is that? The reason why this is so painfully slow, is because every time we generate a batch we read the data in full resolution (1918x1280x3) and resize it. And we do this for every epoch! Therefore, it would make more sense to either store the data in a lower resolution and then to pick the data up, or store the data in cache and access it when it's needed. Or both. Let's slightly change our custom `SegmentationDataSet1` class (create a new class called `SegmentationDataSet2`):

```
1  import torch  
2  from skimage.io import imread  
3  from torch.utils import data  
4  from tqdm import tqdm  
5  
6  
7  class SegmentationDataSet2(data.Dataset):  
8      """Image segmentation dataset with caching and pretransforms."""  
9      def __init__(self,  
10                 inputs: list,  
11                 targets: list,
```

Get started

Open in app



```
15         ):
16         self.inputs = inputs
17         self.targets = targets
18         self.transform = transform
19         self.inputs_dtype = torch.float32
20         self.targets_dtype = torch.long
21         self.use_cache = use_cache
22         self.pre_transform = pre_transform
23
24         if self.use_cache:
25             self.cached_data = []
26
27             progressbar = tqdm(range(len(self.inputs)), desc='Caching')
28             for i, img_name, tar_name in zip(progressbar, self.inputs, self.targets):
29                 img, tar = imread(str(img_name)), imread(str(tar_name))
30                 if self.pre_transform is not None:
31                     img, tar = self.pre_transform(img, tar)
32
33                 self.cached_data.append((img, tar))
34
35         def __len__(self):
36             return len(self.inputs)
37
38         def __getitem__(self,
39                         index: int):
40             if self.use_cache:
41                 x, y = self.cached_data[index]
42             else:
43                 # Select the sample
44                 input_ID = self.inputs[index]
45                 target_ID = self.targets[index]
46
47                 # Load input and target
48                 x, y = imread(str(input_ID)), imread(str(target_ID))
49
50                 # Preprocessing
51                 if self.transform is not None:
52                     x, y = self.transform(x, y)
53
54                 # Typecasting
55                 x, y = torch.from_numpy(x).type(self.inputs_dtype), torch.from_numpy(y).type(self.target
```

[Get started](#)[Open in app](#)

customdatasets2.py hosted with ❤ by GitHub

[view raw](#)

Here we added the argument `use_cache` and `pre_transform`. We basically just iterate over our input and target list and store the images in a list when we instantiate our dataset. When `__getitem__` is called, an image-target pair from this list is returned. I added the `pre_transform` argument because I don't want to change the original files. Instead, I want the images to be picked up, resized and stored in memory. Again, I included a progress bar to visualize the caching. Let's try it out. The changes in code are the following:

```
1  # pre-transformations
2  pre_transforms = ComposeDouble([
3      FunctionWrapperDouble(resize,
4                             input=True,
5                             target=False,
6                             output_shape=(128, 128, 3)),
7      FunctionWrapperDouble(resize,
8                             input=False,
9                             target=True,
10                            output_shape=(128, 128),
11                            order=0,
12                            anti_aliasing=False,
13                            preserve_range=True),
14  ])
15
16 # training transformations and augmentations
17 transforms_training = ComposeDouble([
18     AlbuSeg2d(albumentations.HorizontalFlip(p=0.5)),
19     FunctionWrapperDouble(create_dense_target, input=False, target=True),
20     FunctionWrapperDouble(np.moveaxis, input=True, target=False, source=-1, destination=0),
21     FunctionWrapperDouble(normalize_01)
22 ])
23
24 # validation transformations
25 transforms_validation = ComposeDouble([
26     FunctionWrapperDouble(create_dense_target, input=False, target=True),
27     FunctionWrapperDouble(np.moveaxis, input=True, target=False, source=-1, destination=0),
28     FunctionWrapperDouble(normalize_01)
29 ])
```

Get started

Open in app



```
32 random_seed = 42
33
34 # split dataset into training set and validation set
35 train_size = 0.8 # 80:20 split
36
37 inputs_train, inputs_valid = train_test_split(
38     inputs,
39     random_state=random_seed,
40     train_size=train_size,
41     shuffle=True)
42
43 targets_train, targets_valid = train_test_split(
44     targets,
45     random_state=random_seed,
46     train_size=train_size,
47     shuffle=True)
48
49 # inputs_train, inputs_valid = inputs[:80], inputs[80:]
50 # targets_train, targets_valid = targets[:80], targets[80:]
51
52 # dataset training
53 dataset_train = SegmentationDataSet2(inputs=inputs_train,
54                                     targets=targets_train,
55                                     transform=transforms_training,
56                                     use_cache=True,
57                                     pre_transform=pre_transforms)
58
59 # dataset validation
60 dataset_valid = SegmentationDataSet2(inputs=inputs_valid,
61                                     targets=targets_valid,
62                                     transform=transforms_validation,
63                                     use_cache=True,
64                                     pre_transform=pre_transforms)
```

custom_dataset_2_example.py hosted with ❤ by GitHub

[view raw](#)

And it looks something like this:

```
# dataloader training
dataloader_training = DataLoader(dataset=dataset_train,
                                batch_size=2,
                                shuffle=True)
```

Get started

Open in app



Caching: 100% ██████████ 76/76 [00:26<00:00, 2.83it/s]
 Caching: 5% █████ 1/20 [00:00<00:07, 2.47it/s]

Image by author

The first progress bar represents the training dataloader and the second the validation dataloader. Let's train again for 2 epochs and see how long it'll take.

```
In [4]: # start training
training_losses, validation_losses, learning_rates = trainer.run_trainer()
```

Progress: 100% ██████████ 2/2 [00:02<00:00, 1.08s/it]

Image by author

Training took about 2 seconds only! That's much better. But there is one part we can still improve. Creating the dataset that reads images and stores them in memory takes a bit of time. When you look at the code and the CPU usage, you'll notice that only one core is used. Let's change it in a way, so that all cores are used. Here I use the [multiprocessing](#) library:

```
1 class SegmentationDataSet3(data.Dataset):
2     """Image segmentation dataset with caching, pretransforms and multiprocessing."""
3     def __init__(self,
4                 inputs: list,
5                 targets: list,
6                 transform=None,
7                 use_cache=False,
8                 pre_transform=None,
9                 ):
10         self.inputs = inputs
11         self.targets = targets
12         self.transform = transform
13         self.inputs_dtype = torch.float32
14         self.targets_dtype = torch.long
15         self.use_cache = use_cache
16         self.pre_transform = pre_transform
17
18         if self.use_cache:
19             from multiprocessing import Pool
```

Get started

Open in app



```
22         with pool() as pool:
23             self.cached_data = pool.starmap(self.read_images, zip(inputs, targets, repeat(se
24
25     def __len__(self):
26         return len(self.inputs)
27
28     def __getitem__(self,
29                     index: int):
30         if self.use_cache:
31             x, y = self.cached_data[index]
32         else:
33             # Select the sample
34             input_ID = self.inputs[index]
35             target_ID = self.targets[index]
36
37             # Load input and target
38             x, y = imread(str(input_ID)), imread(str(target_ID))
39
40             # Preprocessing
41             if self.transform is not None:
42                 x, y = self.transform(x, y)
43
44             # Typecasting
45             x, y = torch.from_numpy(x).type(self.inputs_dtype), torch.from_numpy(y).type(self.target
46
47         return x, y
48
49     @staticmethod
50     def read_images(inp, tar, pre_transform):
51         inp, tar = imread(str(inp)), imread(str(tar))
52         if pre_transform:
53             inp, tar = pre_transform(inp, tar)
54         return inp, tar
```

customdataset3.py hosted with ❤ by GitHub

[view raw](#)

Please note that there is no progressbar in this dataset class!

Before we perform training, let's also make a quick detour and talk about the learning rate.

[Get started](#)[Open in app](#)

training. Choosing proper learning rates throughout the learning procedure is difficult as a small learning rate leads to slow convergence while a high learning rate can cause divergence. Also, frequent parameter updates with high variance in SGD can cause fluctuations, which makes finding the (local) minimum for SGD even more difficult. To identify an optimal learning rate, we can test different learning rates empirically with a learning rate range test. Inspired by the best practices I picked up from the fast.ai course, I recommend using a learning rate finder before starting the actual training. Sylvain Gugger from fast.ai wrote a really good [summary](#) about this problem. The code that I will show you is based on Tanjid Hasan Tonmoy's [pytorch-lr-finder](#), which is an implementation of the learning rate range test from Leslie Smith. I only slightly modified the code and included a progressbar (yes, I like them).

```
1  import pandas as pd
2  import torch
3  from torch import nn
4  from matplotlib import pyplot as plt
5  from tqdm import tqdm, trange
6  import math
7
8
9  class LearningRateFinder:
10     """
11     Train a model using different learning rates within a range to find the optimal learning rat
12     """
13
14     def __init__(self,
15                 model: nn.Module,
16                 criterion,
17                 optimizer,
18                 device
19                 ):
20         self.model = model
21         self.criterion = criterion
22         self.optimizer = optimizer
23         self.loss_history = {}
24         self._model_init = model.state_dict()
25         self._opt_init = optimizer.state_dict()
26         self.device = device
```


[Get started](#)[Open in app](#)

```
30         steps=100,
31         min_lr=1e-7,
32         max_lr=1,
33         constant_increment=False
34     ):
35         """
36         Trains the model for number of steps using varied learning rate and store the statistics
37         """
38         self.loss_history = {}
39         self.model.train()
40         current_lr = min_lr
41         steps_counter = 0
42         epochs = math.ceil(steps / len(data_loader))
43
44         progressbar = trange(epochs, desc='Progress')
45         for epoch in progressbar:
46             batch_iter = tqdm(enumerate(data_loader), 'Training', total=len(data_loader),
47                               leave=False)
48
49             for i, (x, y) in batch_iter:
50                 x, y = x.to(self.device), y.to(self.device)
51                 for param_group in self.optimizer.param_groups:
52                     param_group['lr'] = current_lr
53                 self.optimizer.zero_grad()
54                 out = self.model(x)
55                 loss = self.criterion(out, y)
56                 loss.backward()
57                 self.optimizer.step()
58                 self.loss_history[current_lr] = loss.item()
59
60                 steps_counter += 1
61                 if steps_counter > steps:
62                     break
63
64                 if constant_increment:
65                     current_lr += (max_lr - min_lr) / steps
66                 else:
67                     current_lr = current_lr * (max_lr / min_lr) ** (1 / steps)
68
69         def plot(self,
70                 smoothing=True,
71                 clipping=True,
```

Get started

Open in app



```

75     Shows loss vs learning rate(log scale) in a matplotlib plot
76     """
77     loss_data = pd.Series(list(self.loss_history.values()))
78     lr_list = list(self.loss_history.keys())
79     if smoothing:
80         loss_data = loss_data.ewm(alpha=smoothing_factor).mean()
81         loss_data = loss_data.divide(pd.Series(
82             [1 - (1.0 - smoothing_factor) ** i for i in range(1, loss_data.shape[0] + 1)]))
83     if clipping:
84         loss_data = loss_data[10:-5]
85         lr_list = lr_list[10:-5]
86     plt.plot(lr_list, loss_data)
87     plt.xscale('log')
88     plt.title('Loss vs Learning rate')
89     plt.xlabel('Learning rate (log scale)')
90     plt.ylabel('Loss (exponential moving average)')
91     plt.show()
92
93     def reset(self):
94         """
95         Resets the model and optimizer to its initial state
96         """
97         self.model.load_state_dict(self._model_init)
98         self.optimizer.load_state_dict(self._opt_init)
99         print('Model and optimizer in initial state.')

```

lr_rate_finder.py hosted with ❤ by GitHub

view raw


Let's perform such a learning rate range test. Since our dataset is rather small (96 images), we'll perform some extra steps (1000). The upper progressbar displays the number of epochs and the lower progressbar shows the number of steps we perform on the current epoch.

```

In [9]: from lr_rate_finder import LearningRateFinder
lr_f = LearningRateFinder(model, criterion, optimizer, device)
lr_f.fit(dataloader_training, steps=1000)

```

Progress: 100%  27/27 [00:17<00:00, 1.52it/s]

Training: 32%  12/38 [00:00<00:00, 49.14it/s]

Get started

Open in app



Let's plot the results of the test:

In [10]: `lrf.plot()`

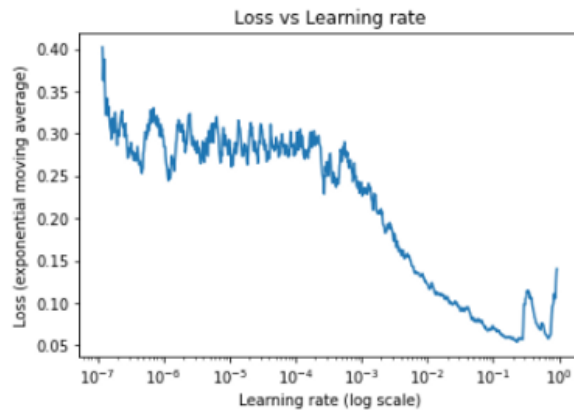


Image by author

0.01 seems to be a good learning rate. We'll take it. Let's train for 100 epochs...

```
In [12]: # trainer
trainer = Trainer(model=model,
                  device=device,
                  criterion=criterion,
                  optimizer=optimizer,
                  training_DataLoader=dataloader_training,
                  validation_DataLoader=dataloader_validation,
                  lr_scheduler=None,
                  epochs=100,
                  epoch=0)

# start training
training_losses, validation_losses, learning_rates = trainer.run_trainer()

Progress: 100%  100/100 [12:09<00:00, 7.30s/it]
```

Image by author

...and visualize the training and validation loss. For that I will use matplotlib and write a function that I can add to the `visual.py` file.

```
1 def plot_training(training_losses,
2                   validation_losses,
3                   learning_rate,
4                   gaussian=True,
5                   sigma=2,
```

Get started

Open in app



```
9     Returns a loss plot with training loss, validation loss and learning rate.
10     """
11
12     import matplotlib.pyplot as plt
13     from matplotlib import gridspec
14     from scipy.ndimage import gaussian_filter
15
16     list_len = len(training_losses)
17     x_range = list(range(1, list_len + 1)) # number of x values
18
19     fig = plt.figure(figsize=figsize)
20     grid = gridspec.GridSpec(ncols=2, nrows=1, figure=fig)
21
22     subfig1 = fig.add_subplot(grid[0, 0])
23     subfig2 = fig.add_subplot(grid[0, 1])
24
25     subfigures = fig.get_axes()
26
27     for i, subfig in enumerate(subfigures, start=1):
28         subfig.spines['top'].set_visible(False)
29         subfig.spines['right'].set_visible(False)
30
31     if gaussian:
32         training_losses_gauss = gaussian_filter(training_losses, sigma=sigma)
33         validation_losses_gauss = gaussian_filter(validation_losses, sigma=sigma)
34
35         linestyle_original = '.'
36         color_original_train = 'lightcoral'
37         color_original_valid = 'lightgreen'
38         color_smooth_train = 'red'
39         color_smooth_valid = 'green'
40         alpha = 0.25
41     else:
42         linestyle_original = '-'
43         color_original_train = 'red'
44         color_original_valid = 'green'
45         alpha = 1.0
46
47     # Subfig 1
48     subfig1.plot(x_range, training_losses, linestyle_original, color=color_original_train, label
49                 alpha=alpha)
```

Get started

Open in app



```

53     subfig1.plot(x_range, training_losses_gauss, '-', color=color_smooth_train, label='Train
54     subfig1.plot(x_range, validation_losses_gauss, '-', color=color_smooth_valid, label='Val
55     subfig1.title.set_text('Training & validation loss')
56     subfig1.set_xlabel('Epoch')
57     subfig1.set_ylabel('Loss')
58
59     subfig1.legend(loc='upper right')
60
61     # Subfig 2
62     subfig2.plot(x_range, learning_rate, color='black')
63     subfig2.title.set_text('Learning rate')
64     subfig2.set_xlabel('Epoch')
65     subfig2.set_ylabel('LR')
66
67     return fig

```

plot_training_loss.py hosted with ❤ by GitHub

view raw

Let's see what the function `plot_training()` will output when we pass in our losses and the learning rate.

```

In [40]: from visual import plot_training
fig = plot_training(training_losses,
                    validation_losses,
                    learning_rates,
                    gaussian=True,
                    sigma=1,
                    figsize=(10, 4))

```

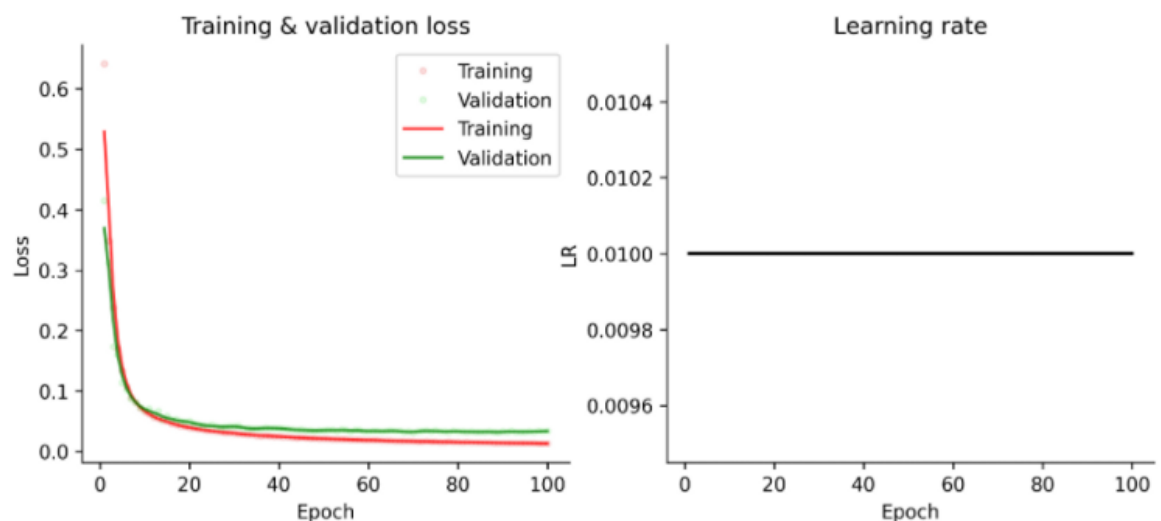


Image by author

[Get started](#)[Open in app](#)

we can then save our model with PyTorch.

```
# save the model
model_name = 'carvana_model.pt'
torch.save(model.state_dict(), pathlib.Path.cwd() / model_name)
```

Summary

In this part, we performed training with a sample of the Carvana dataset by creating a simple training loop. The progress of this training loop can be visualized with a progressbar and the result of training can be plotted with matplotlib. We noticed that training was painfully slow because our data was picked up very slowly by our custom data generator. Because of that, we changed it in a way so that data is only read once and then picked up from memory when needed. We also made use of multiprocessing for that case. Additionally, we added a learning rate range finder, to determine an optimal learning rate which we then used for model training.

In the [next chapter](#), we'll let the model predict the segmentation maps of unseen image data (inference).

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

[Get this newsletter](#)

You'll need to sign in or create an account to receive this newsletter.

[Python](#)[Deep Learning](#)[Semantic Segmentation](#)[Pytorch](#)[Tutorial](#)

[Get started](#)[Open in app](#)

Get the Medium app

