Google Developer Groups
Cloud • Pune

Google Cloud

Community Day 2023

Google Cloud

Community
Day 2023

Introduction to Graph Data Science

Dr. Yogesh Haribhau Kulkarni

AI Coach

Google Developer Groups
Cloud • Pune

Google Developer Groups
Cloud • Pune

Hello!

I am Yogesh Kulkarni

- 20+ years in CAD/Engineering software development
- Got Bachelors, Masters and Doctoral degrees in Mechanical Engineering (specialization: Geometric Modeling Algorithms).
- Currently doing Coaching in fields such as Data Science, Artificial Intelligence Machine-Deep Learning (ML/DL) and Natural Language Processing (NLP).

Feel free to follow me at:

- LinkedIn (www.linkedin.com/in/yogeshkulkarni/)
- Medium (yogeshharibhaukulkarni.medium.com)

1.

What are Graphs?

First, one disclaimer

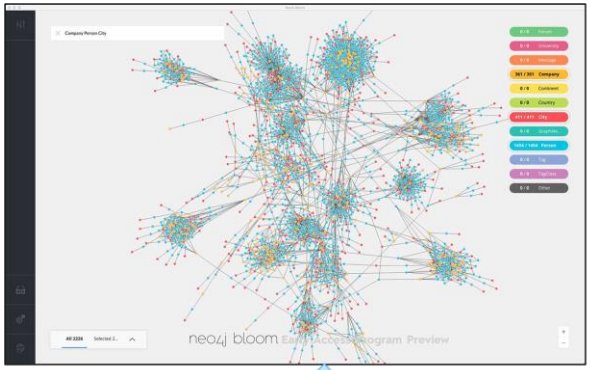

This is a graph

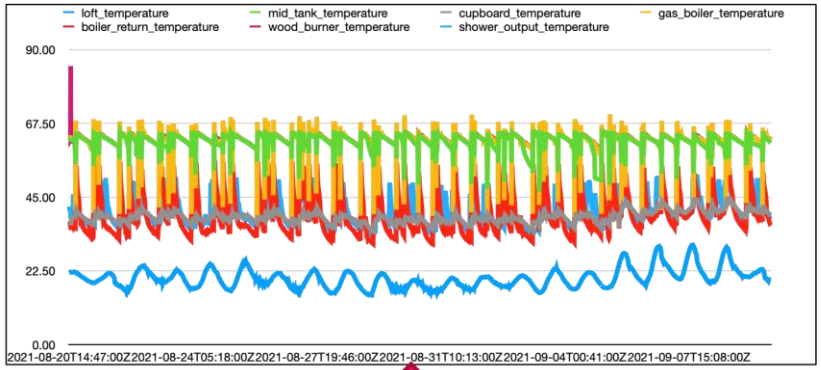

This is a chart

Everything is Naturally Connected

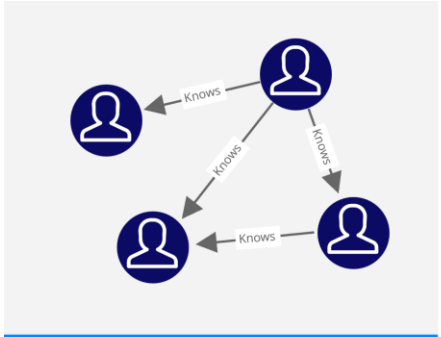

Networks of People

Employees, Customers, Suppliers, Partners, Influencers, etc.

Transaction Networks

Risk management, Supply chain, Orders, Payments, etc.

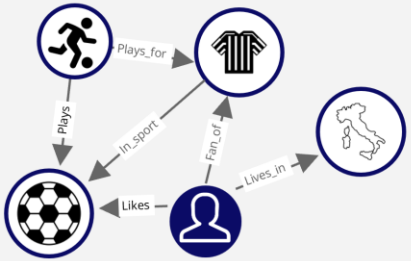

Knowledge Networks

Enterprise content, Domain specific content, eCommerce content, etc

Take advantage of

- graphs to capture the network structure
- graph queries to store and retrieve relationships
- graph algorithms to infer relationships

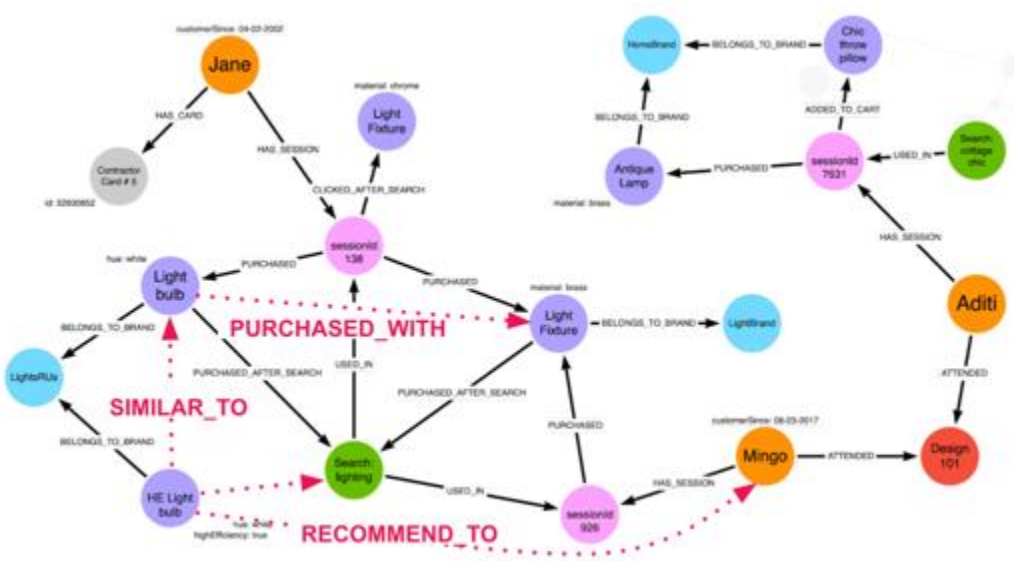

What is data science?

"Data science is an interdisciplinary field that uses scientific methods, processes, algorithms and systems to extract knowledge and insights from structured and unstructured data." - Wikipedia

Mark Needham - Intro to Graph Data Science with Neo4j

What is GRAPH data science?

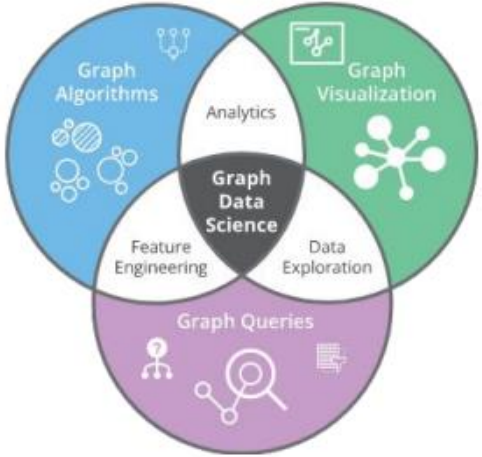

Graph Data Science is a science-driven approach to gain knowledge from the relationships and structures in data, typically to power predictions.

Data scientists use relationships to answer questions.

What's different?

- Separate vs related
- Fixed vs variable size

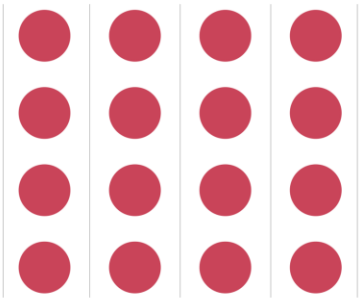

Static approach

Discrete data, as represented in a relational database

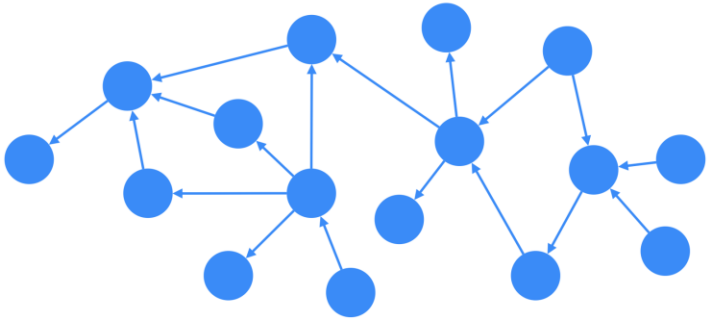

Connected approach, or "graph"

Data represented as network in a "graph"-structure

Graphs enrich all stages of Data Science

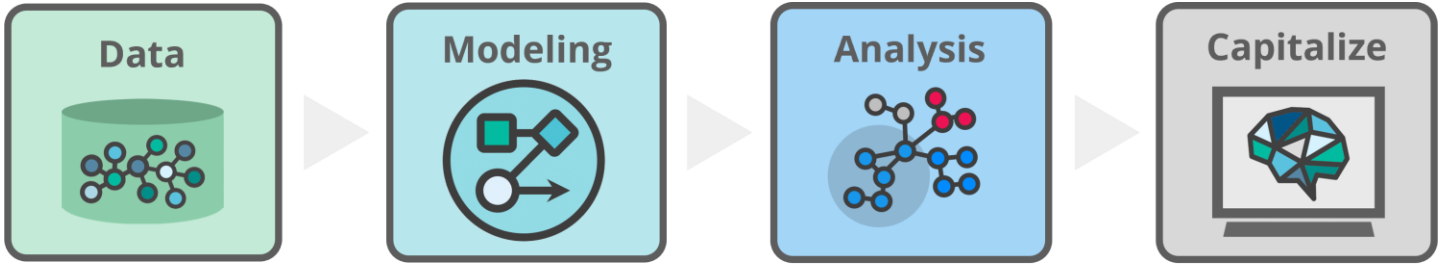

Knowledge Graphs

Graph Feature Engineering and Graph ML

Graph Analytics, Investigations and Counterfactuals

Integrations and Knowledge Graphs for Heuristic AI

Graph Data Science Applications

Fraud Detection

Disambiguation & Segmentation

Personalized Recommendations

Life Sciences

Churn Prediction

Search & Master Data Mgmt.

Predictive Maintenance

Cybersecurity

Section 1/3

Knowledge Graphs

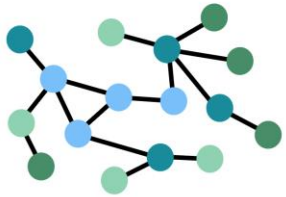

Find the patterns you're looking for in connected data

Graph Algorithms

Use unsupervised machine learning techniques to identify associations.

Graph Native Machine Learning

Use embeddings to learn the features in your graph that you don't even know are important yet.

Train supervise ML models to predict links, labels, and missing data.

Knowledge Graph - Building Blocks

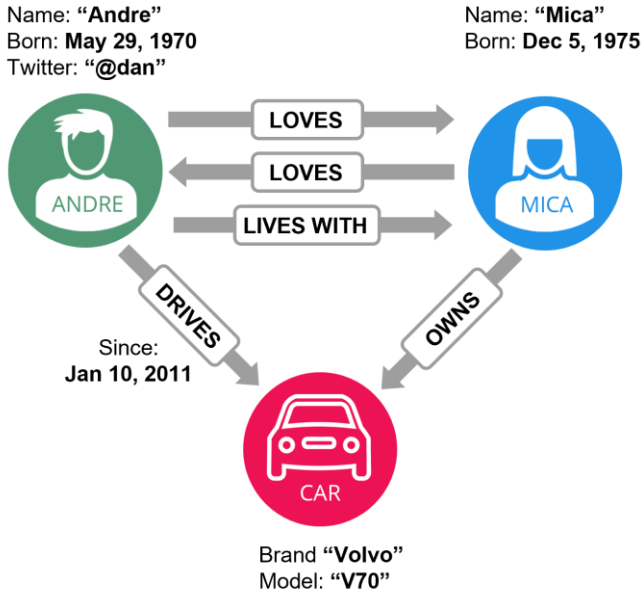

Node

Represents an entity in the graph

Relationship

Connect nodes to each other

Property

Describes a node or relationship: e.g. name, age, weight etc

Database, Query Language and Visualization

DATA

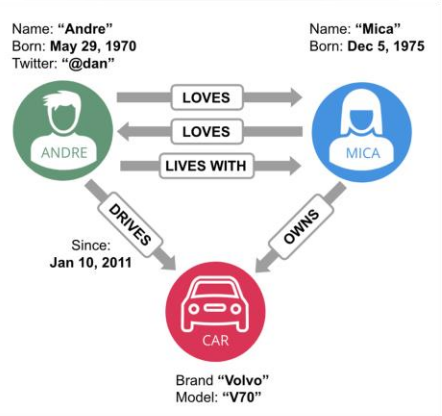

GRAPH QUERIES

```
MATCH (a)-[:KNOWS]->(b)
RETURN b.name
UNION ALL
MATCH (a)-[:LOVES]->(b)
RETURN b.name
```

CYPHER

https://neo4j.com/developer/cypher/

GRAPH VISUALIZATION

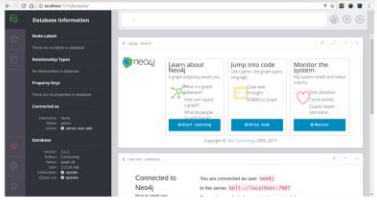

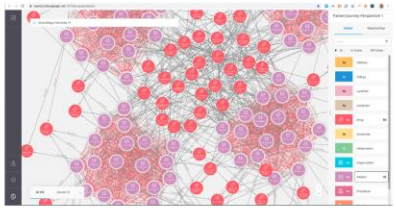

What can you do with a knowledge graph?

Financial Domain

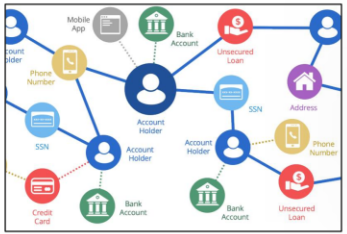

How many flagged accounts are in the applicant's network 4+ hops out?

How many login / account variables in common?

Add these metrics to your approval process

Life Sciences

What completes the connections from genes to diseases to targets?

What genes can be reached 4+ hops out from a known drug target?

What mechanisms in common are there between two drugs?

Marketing and Recommendations

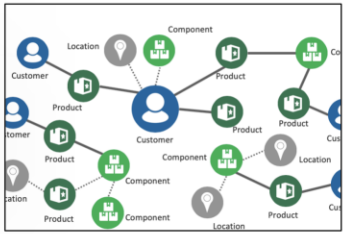

Collaborative filtering: users who bought X, also bought Y (open-ended pattern matching)

What items make you more likely to buy additional items in subsequent transactions?

Traverse hierarchies - what items are similar 4+ hops out?

Section 2/3

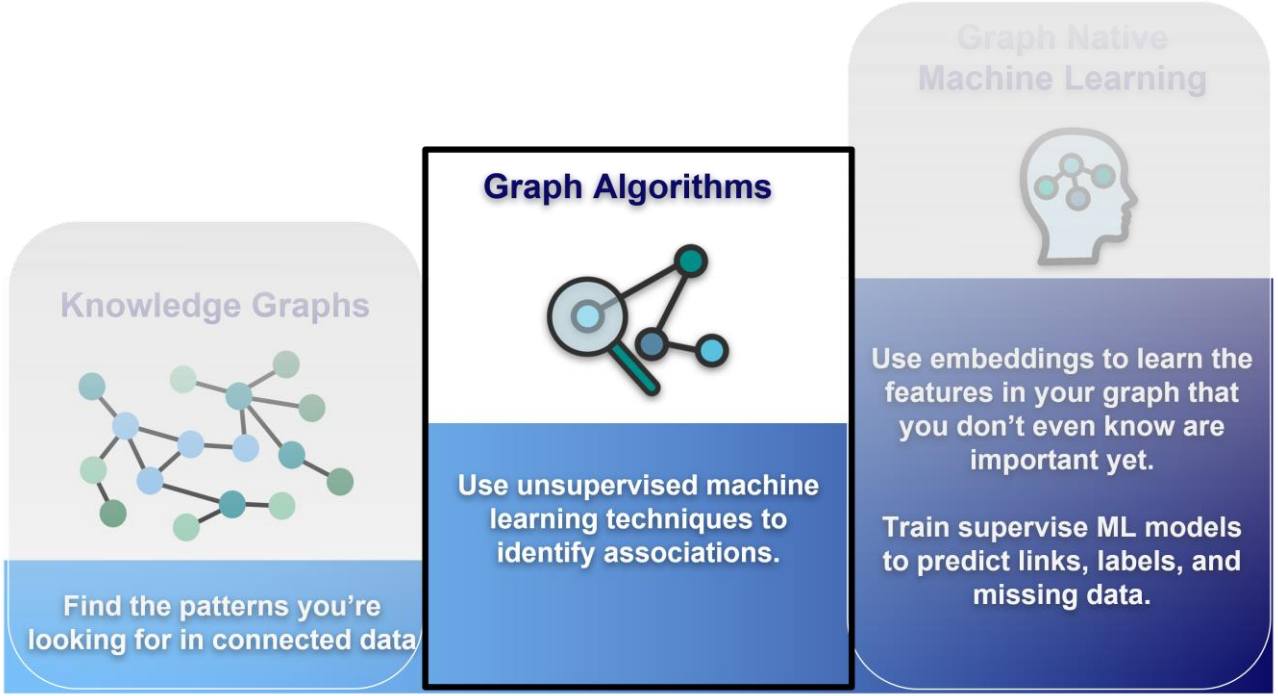
Knowledge Graphs
Find the patterns you're looking for in connected data
Graph Algorithms
Use unsupervised machine learning techniques to identify associations.
Graph Native Machine Learning
Use embeddings to learn the features in your graph that you don't even know are important yet.
Train supervise ML models to predict links, labels, and missing data.

Graph Algorithm types

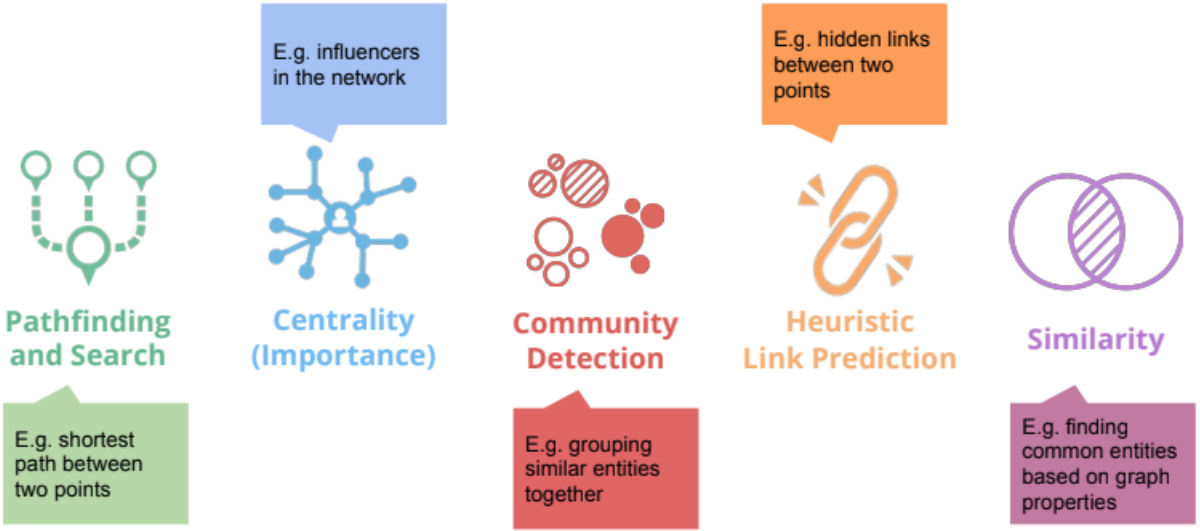

Mark Needham - Intro to Graph Data Science with Neo4j

Graph Algorithms

Pathfinding & Search

- Shortest Path
- Single-Source Shortest Path
- All Pairs Shortest Path
- A* Shortest Path
- Yen's K Shortest Path
- Minimum Weight Spanning Tree
- K-Spanning Tree (MST)
- Random Walk
- Breadth & Depth First Search

Centrality & Importance

- Degree Centrality
- Closeness Centrality
- Harmonic Centrality
- Betweenness Centrality & Approx.
- PageRank
- Personalized PageRank
- ArticleRank
- Eigenvector Centrality
- Hyperlink Induced Topic Search (HITS)
- Influence Maximization (Greedy, CELF)

Community Detection

- Triangle Count
- K-Means
- Local Clustering Coefficient
- Connected Components (Union Find)
- Strongly Connected Components
- Label Propagation
- Louvain Modularity
- K-1 Coloring
- Modularity Optimization
- Speaker Listener Label Propagation

Supervised Machine Learning

- Node Classification
- Link Prediction
- Node Regression

Heuristic Link Prediction

- Adamic Adar
- Common Neighbors
- Preferential Attachment
- Resource Allocations
- Same Community
- Total Neighbors

Similarity

- Node Similarity
- K-Nearest Neighbors (KNN)
- Jaccard Similarity
- Cosine Similarity
- Pearson Similarity
- Euclidean Distance
- Approximate Nearest Neighbors (ANN)

Graph Embeddings

- Node2Vec
- FastRP
- GraphSAGE

… and more!

- Synthetic Graph Generation
- Scale Properties
- Collapse Paths
- One Hot Encoding
- Split Relationships
- Graph Export
- Pregel API (write your own algos)

What are Graph Algorithms?

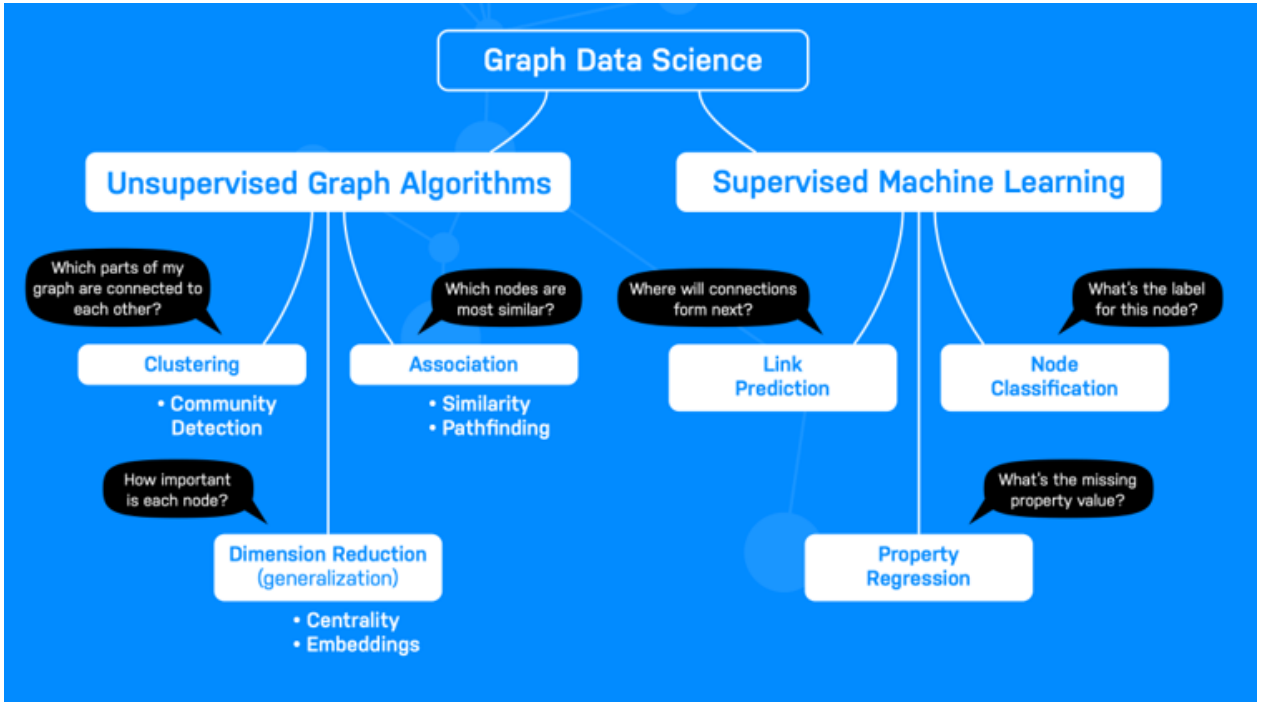

Section 3/3

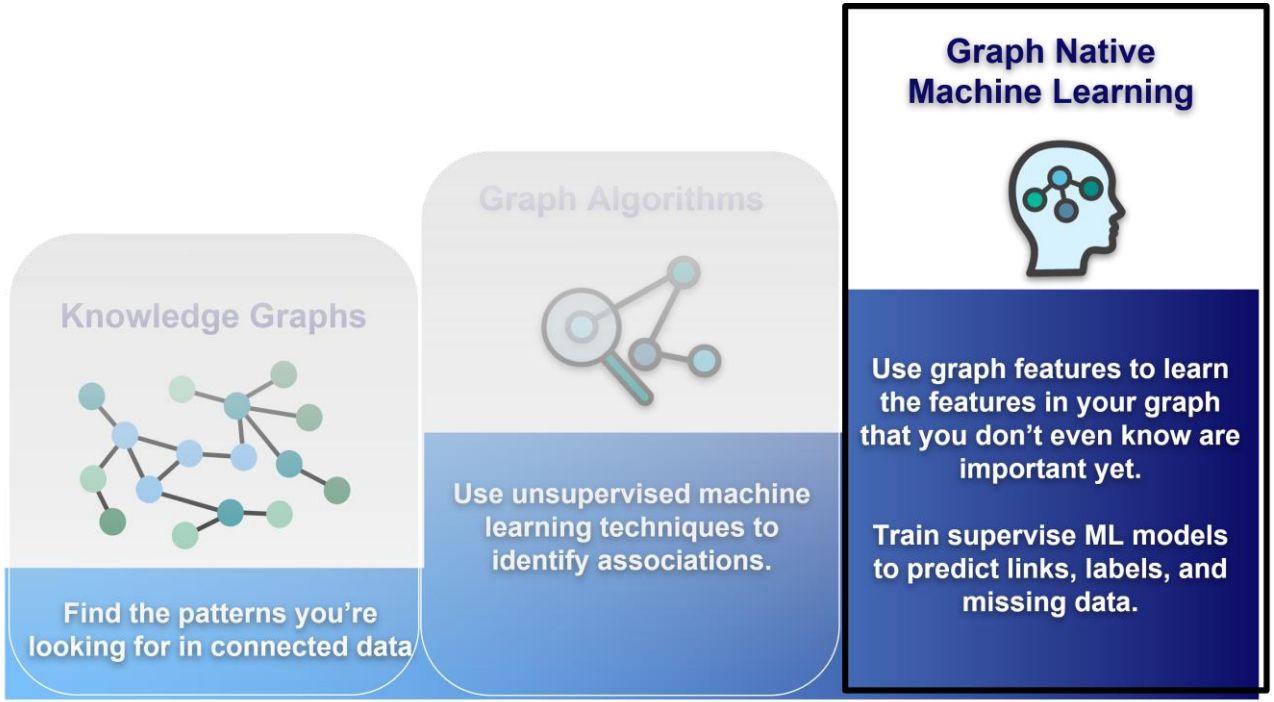
Knowledge Graphs
Find the patterns you're looking for in connected data
Graph Algorithms
Use unsupervised machine learning techniques to identify associations.
Graph Native Machine Learning
Use graph features to learn the features in your graph that you don't even know are important yet.
Train supervise ML models to predict links, labels, and missing data.

Relationship-Driven AI

- Traditional ML ignore network structure because it's difficult to extract
- Use the right data structures to store and retrieve relationships
- Add relationships to AI/ML pipelines to make them contextual and to unlock otherwise unattainable predictions

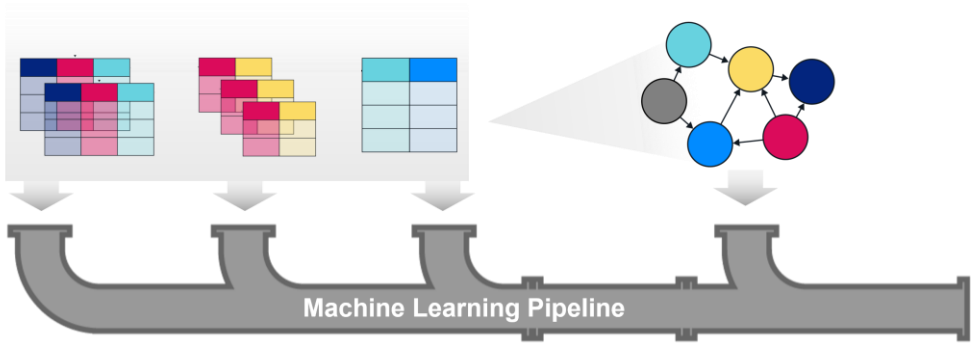

Graph Feature Engineering

Queries

Human-crafted query, human-readable result

```
MATCH (p1:Person)-[:ENEMY]->(:Person)<-[:ENEMY]-(p2:PERSON)
MERGE (p1)-[:FRIEND]->(p2)
```

Algorithms

Predefined formula, human-readable result

$$PR(u) = \sum_{v \in B_u} \frac{PR(v)}{L(v)}$$

```
PageRank(Emil)    = 13.25
PageRank(Amy)     =  4.83
PageRank(Alicia)  =  4.75
```

Embeddings

AI-learned formula, machine-readable result

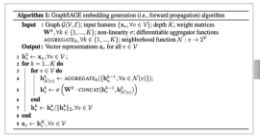

```
Node2Vec(Emil)    =[5.4 5.1 2.4 4.5 3.1]
Node2Vec(Amy)     =[2.8 1.8 7.2 0.9 3.0]
Node2Vec(Alicia)=[1.4 5.2 4.4 3.9 3.2]
```

Machine Learning Workflows

Train ML models based on results

$$f(x)$$

Graph Machine Learning

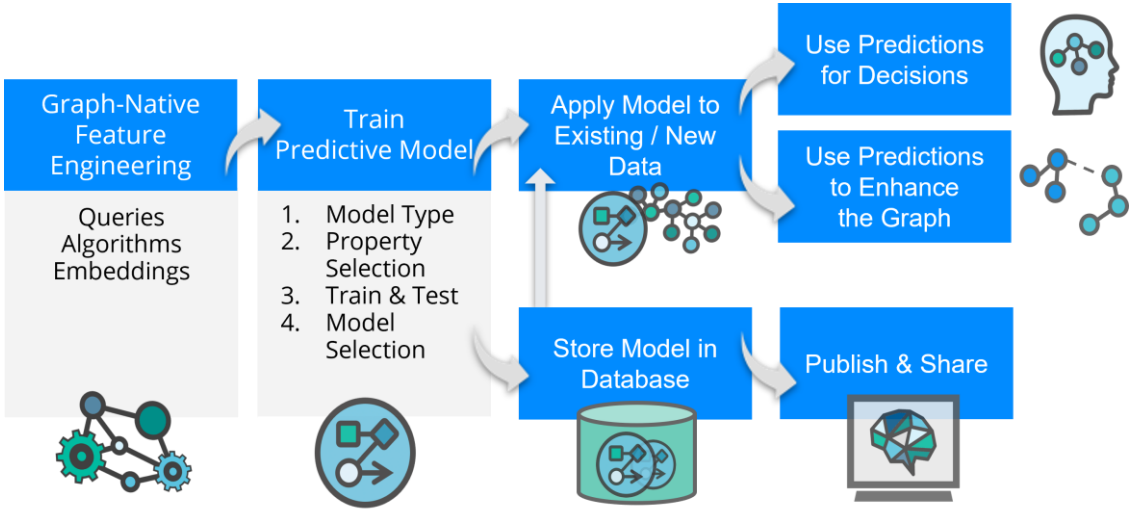

The Only Completely In-Graph, ML Workflow

In-Graph Machine Learning

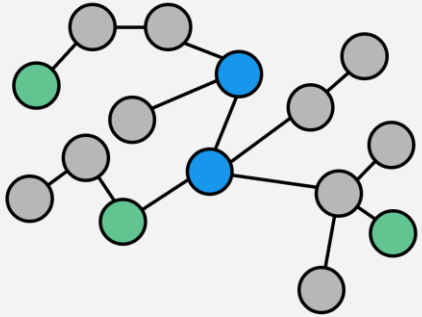

Labeled data: Pairs of nodes that are either linked or not
Features: Pre-existing attributes, algorithms (pageRank), embedding

$f(x)$

Node classification: "What kind of node is this?"

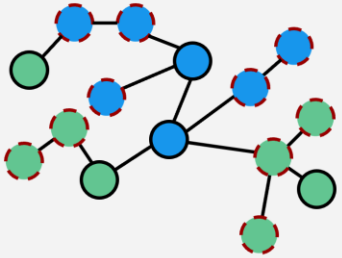

$f(x)$

Link prediction: "Should there be a relationship between these nodes?"

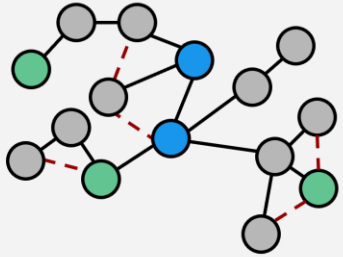

Neo4j Graph Data Science Framework

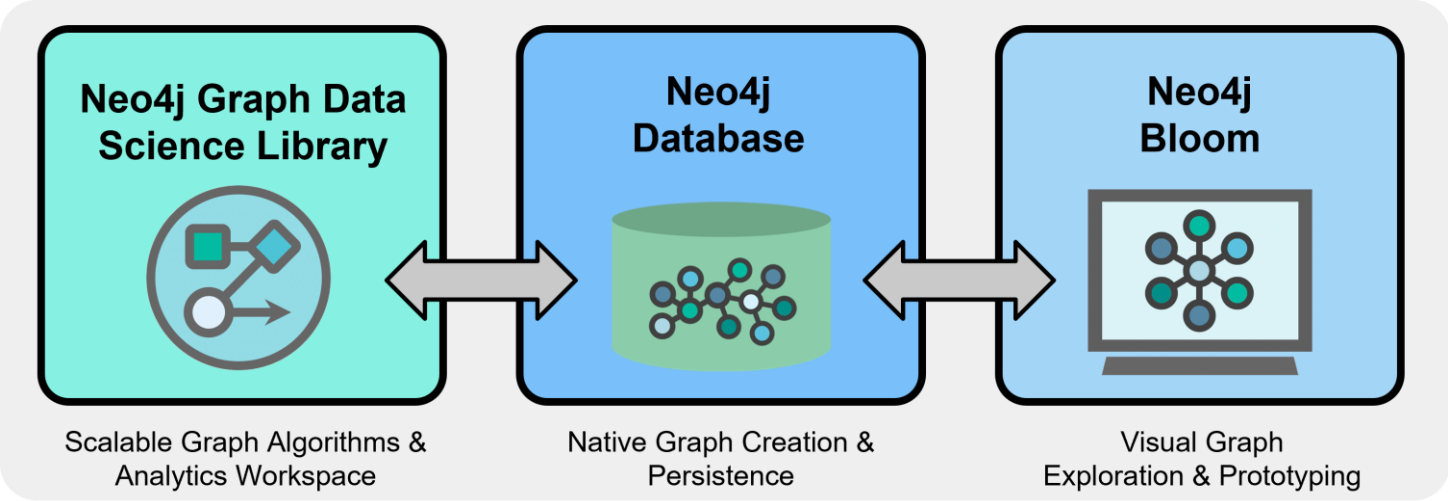

What's available in GDS library?

Pathfinding & Search

- Breadth/Depth First Search
- Shortest Path
- Single-Source Shortest Path
- All Pairs Shortest Path
- Minimum Spanning Tree
- A* Shortest Path
- Yen's K Shortest Path
- K-Spanning Tree (MST)
- Random Walk

Centrality / Importance

- Degree Centrality
- Closeness Centrality
- CC Variations: Harmonic, Dangalchev, Wasserman & Faust
- **Betweenness Centrality**
- **PageRank**
- Personalized PageRank
- ArticleRank
- Eigenvector Centrality

Community Detection

- **Triangle Count**
- **Clustering Coefficients**
- **Weakly Connected Components**
- Strongly Connected Components
- **Label Propagation**
- **Louvain Modularity**
- K1 coloring
- Modularity optimization

Similarity

- **Node Similarity**
- K-Nearest Neighbors
- Euclidean Distance
- Cosine Similarity
- Jaccard Similarity
- Overlap Similarity
- Pearson Similarity

Graph Embeddings

- **Node Similarity**
- K-Nearest Neighbors
- Euclidean Distance
- Cosine Similarity
- Jaccard Similarity
- Overlap Similarity
- Pearson Similarity

Link Prediction

- Adamic Adar
- Common Neighbors
- Preferential Attachment
- Resource Allocations
- Same Community
- Total Neighbors

...and also:

- Random graph generation
- One hot encoding
- Machine learning models
  - Node Classification
  - Link Prediction

Mark Needham - Intro to Graph Data Science with Neo4j

GDS syntax

```
CALL gds[.<tier>].<algorithm>.<execution-mode>[.<estimate>](
    graphName: STRING,
    configuration: MAP
)
```

<tier> - tier of support

<algorithm> - name of algorithm to execute

<execution-mode> - how to run the algorithm: write/stream/stats

<estimate> - estimate the memory requirements of the algorithm

Mark Needham - Intro to Graph Data Science with Neo4j

Where do I find the GDS library?

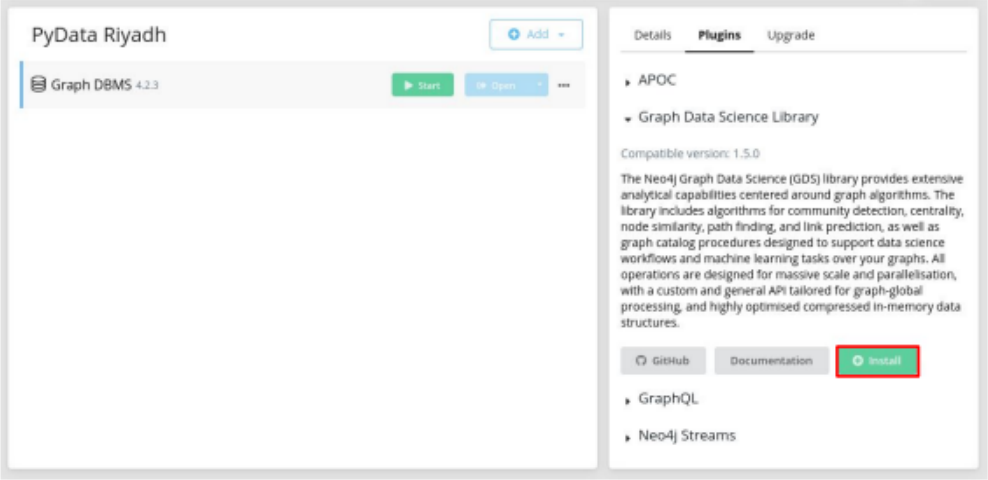

neo4j.com/developer/neo4j-desktop/

Mark Needham - Intro to Graph Data Science with Neo4j

Where do I find the GDS library?

- Download from: neo4j.com/download-center
- Documentation: neo4j.com/docs/graph-data-science/1.5

What Next?

Free online training and certification

- *dev.neo4j.com/learn*
- *dev.neo4j.com/datasets*

connect with the community:

- *dev.neo4j.com/chat*
- *dev.neo4j.com/forum*
- *dev.neo4j.com/newsletter*

Next developer events

- *Live Streams - Weekly & Online*
- *Local Meetups neo4j.com/events*

Get started - downloads & GitHub projects

- *dev.neo4j.com/sandbox*
- *github.com/neo4j*

Google Developer Groups
Cloud • Pune

Continue your graph journey
with Graph Academy

What is Graph Academy?
Free, Self-Paced, Hands-on Online Training to help you learn how to build, optimize and launch your Neo4j project, all from the Neo4j experts.

What's more?
2 free certifications designed to test you on your overall knowledge of Neo4j:

- Neo4j Graph Data Science Certification
- Neo4j Certified Professional

Interested? For more information visit:
www.graphacademy.neo4j.com

neo4j | NODES 23

Beginner to Expert. It All Connects Here.

October 26, 2023

Save the Date

Submit your talk by June 30, 2023

dev.neo4j.com/nodes23

References

- Neo4j Graph Data Science - Alison Cossette
- Neo4j supplied presentations: Graph Databases, Cypher.
- “an introduction to neo4j (graph database tutorial for beginners)" - Chris Hay
- CIS 6930 - Advanced Databases - Neo4j
- Learning Neo4j" - Wabri/LearningNeo4j - Github
- Neo4J Certification Sample Questions - https://wiki.glitchdata.com/index.php/Neo4J_Certification
- Neo4j Certified Professional: Exam Practice Tests - By Cristian Scutaru
- "Graph Data Science with Neo4j Graph Algorithms - Will Lyon" Youtube
- "Mark Needham - Intro to Graph Data Science with Neo4j" Youtube
- "THE POWER OF “GRAPH DATA SCIENCE” using Neo4j" Youtube
- "Using Neo4j Graph Data Science in Python to Improve Machine Learning Models" - Tomaz Bratanic
- "Graph Data Science for Supply Chains – Part 1" - Zach Blumenfeld

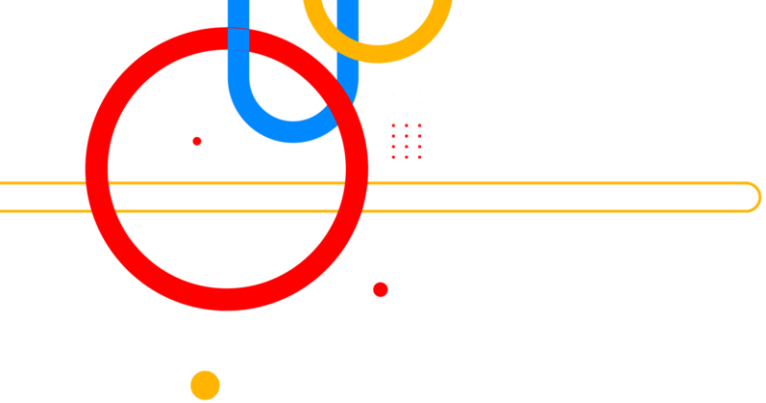

Google Developer Groups
Cloud • Pune

Thank You !

yogeshkulkarni@yahoo.com